

# I can't believe there's no images!

## Learning Visual Tasks Using Only Language Supervision

Sophia Gu\* Christopher Clark\* Aniruddha Kembhavi

Allen Institute for Artificial Intelligence

{sophiag, chrisc, anik}@allenai.org

[IMP]: goal of the paper is to transfer info/style from the textual data to vision tasks without training the visual on visual training data.

### Abstract

Many high-level skills that are required for computer vision tasks, such as parsing questions, comparing and contrasting semantics, and writing descriptions, are also required in other domains such as natural language processing. In this paper, we ask whether it is possible to learn those skills from text data and then transfer them to vision tasks without ever training on visual training data. Key to our approach is exploiting the joint embedding space of contrastively trained vision and language encoders. In practice, there can be systematic differences between embedding spaces for different modalities in contrastive models, and we analyze how these differences affect our approach and study strategies to mitigate this concern. We produce models using only text training data on four representative tasks: image captioning, visual entailment, visual question answering and visual news captioning, and evaluate them on standard benchmarks using images. We find these models perform close to models trained on images, while surpassing prior work for captioning and visual entailment in this text-only setting by over 9 points, and outperforming all prior work on visual news by over 30 points. We also showcase a variety of stylistic image captioning models that are trained using no image data and no human-curated language data, but instead using readily-available text data from books, the web, or language models.

### 1. Introduction

Although vision and natural language processing (NLP) tasks are typically thought of as being very distinct, there is often a high degree of overlap in the skills needed to complete them. Visual question answering and reading comprehension question answering both require parsing and understanding questions, visual entailment and textual entailment require comparing different semantic meanings, and captioning and summarization require writing text that sum-

marizes the semantics of the input. This raises an intriguing possibility: if a model learned to complete one of these tasks using a high-level semantic representation of the input text, then in theory it could immediately be able to complete the corresponding visual task as long as the input image is encoded in the same semantic representation. We call this challenge zero-shot cross-modal transfer because it requires applying skills learned from one modality to a different one. Achieving this would be a step towards building multi-modal models that can generalize skills across modalities without needing expensive training data for each modality, and has potential applications for tasks where visual training data is scarce but text data is relatively easy to collect.

Accomplishing this requires encoding images and text into a shared semantic space. We use vision and language (V&L) models trained with a contrastive loss for this purpose [51, 25]. These models learn to embed text and images into vectors such that the vectors for matching images and captions are close together, and vectors for unrelated images and captions are far apart. Although this loss was originally intended for representation learning and zero-shot classification, here we show it also facilitates cross-modal transfer.

To do this, we propose a method called Cross modal transfer On Semantic Embeddings (CLOSE). An outline of CLOSE is shown in Figure 1. During training, the text inputs are encoded into a vector using the (frozen) text encoder from a contrastive model, which is then used as an input to a model. During testing, the visual input is embedded with a (frozen) image encoder and used in place of the text embedding. Because these encoders were explicitly trained to produce embeddings that encode semantics in similar ways, learning to read and process the text vector should naturally translate to the ability to read and process the image vector. Although we focus on text-to-image transfer in this paper, our approach is applicable to other contrastive models such as videos [75], point clouds [1], and audio [22, 11, 73], potentially allowing transfer between many other modalities.

One potential difficulty with this approach is that, while contrastive embeddings do share some structure between

\*Equal contribution

use of this technique

Limitation

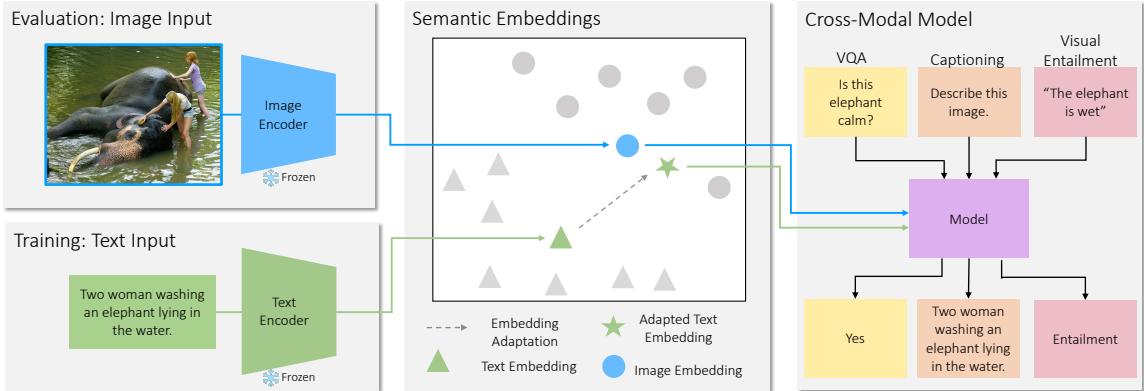


Figure 1: Overview of CLOSE. During training, input text is encoded into a vector with a text encoder and adapted with an adaptation method. A model learns to use the vector to perform a task such as VQA, captioning, or visual entailment. During testing, an input image is encoded with an image encoder instead to allow cross-modal transfer.

modalities, there can still be significant differences between the image and text vectors in practice [39]. To mitigate this, we propose to additionally use *adapters* that modify the text vectors being used during training. We find adding Gaussian noise to be very effective in boosting performance, but consider other approaches as well in our analyses.

Text-to-image transfer is a relatively unexplored setting, so we first conduct extensive experiments to establish that CLOSE can handle the text-to-image domain shift without a major performance drop. We compare models trained with CLOSE on text alone to models trained with images and text on three standard V&L tasks: captioning, visual questioning answers (VQA) and visual entailment, and the more complex task of visual news captioning [40]. We find the text-only models generally perform reasonably close to versions trained with images, showing that CLOSE can effectively transfer many skills across modalities. We surpass the previous best text-only method in captioning [79] by 17 CIDEr (78.2 vs. 95.4) and visual entailment [57] by 9 points (66.6 vs. 75.9), making our method state-of-the-art for these settings by a large margin. There are no prior results for VQA and visual news in this setting, however we do surpass the previously best reported result in visual news even with images [40] (50.5 vs 80.8 CIDEr).

These experiments show that efficient text-to-image transfer is possible. This has important practical implications because text training data can be directly constructed by annotators, mined from many existing text datasets, or even generated by a large language model such as GPT-3 [4], and can therefore be significantly less expensive than constructing visual training data. We demonstrate this potential by training effective CLOSE captioning models from text generated by large language models [4], meaning the only human annotation required was for prompt construction. We also train several stylistic captioning models without any labeled images (see Figure 2). We collect text

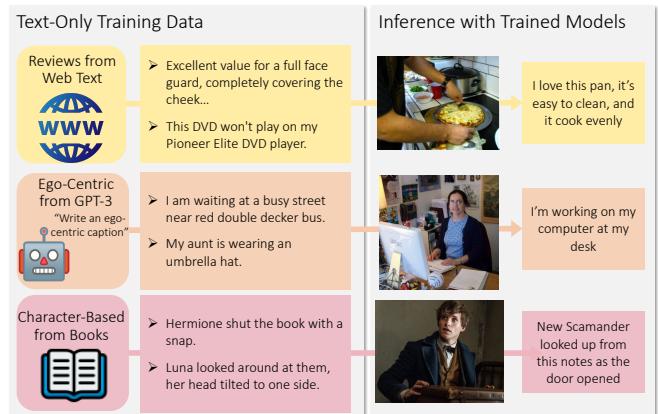


Figure 2: Using CLOSE to learn stylistic captioning without image data. Text examples of the desired style are gathered from sources such as the web, books, or GPT-3. Models are trained on text only and then applied to images.

with various styles from a diverse set of sources, including internet reviews, books, and GPT-3 generations, and demonstrate that CLOSE models trained on this text can produce accurate and stylistically correct captions for images.

Finally, we complete two analyses: A sensitivity analysis showing that CLOSE is robust to cases where text and image vectors differ by a constant offset, which therefore allows CLOSE to work despite seemingly large differences between the image/text embeddings. Additionally, a study on the effectiveness of using an auxiliary vision and language corpus to build an improved adapter. We find that improvements are possible but vary depending on the source of that data and that a particularly effective approach is to use the auxiliary data to compute a structured covariance matrix for use when adding Gaussian noise.

In summary, our contributions include: (i) introducing

adapter is used to close the modality gap and it is done by using auxiliary vision-lang data to compute the noise in the model

the CLOSE model for zero-shot cross-modal transfer; (ii) showing that training CLOSE with text data alone, on four V&L tasks, gives results close to models trained on both images and text; (iii) SoTA results when using only text for three of the tasks; (iv) demonstrating an application of CLOSE for stylistic captioning; (v) analyzing how differences between image/text vectors in contrastive models and how different adapters affect CLOSE’s performance. To facilitate future work in the community, we release our code<sup>1</sup>.

## 2. Method

**Model.** Our approach uses the image/text encoder from a contrastive model to encode the input, and then follows many prior works (e.g., [27, 7]) by fine-tuning a pre-trained language model to process the input vector, along with any additional input text, to generate output text. First, the input image or text vector is normalized to have unit length to match what is used in the contrastive loss. Then that vector is converted into a number of vectors, we use 4 in our experiments, of the same dimensionality as the language model’s embedding layer using a linear layer. Next, other input text (e.g., the hypothesis in visual entailment or the question in VQA) is tokenized and embedded with the language model’s embedding layer. Those embeddings are concatenated with the embeddings built from the input vector to construct an input sequence for the language model.

For the sake of simplicity, we train the model generatively for all tasks [20, 8]. The model generates a caption, a free-form question answer, or a class name for the tasks of captioning, VQA, and visual entailment respectively. During training, the language model and linear layer are fine-tuned, but the text encoder is kept frozen to ensure the correspondence between text and image vectors learned during pre-training is preserved.

**Modality Gap.** In practice, text and image vectors from contrastive models can be far apart, a phenomenon known as the modality gap [39]. For example, on COCO captions [6] the average cosine similarity between an image and paired caption is only 0.26, while the average similarity between two unrelated captions is 0.35. Figure 3a shows this gap causes image and text vectors to fall into separate clusters in the vector space. The root cause is that the cross-entropy loss used by contrastive models only requires paired image and text vectors to be close relative to random image and text pairs, which does not necessarily mean they are close in absolute terms, see Liang *et al.* [39] for more discussion.

We thus adopt a simple and effective solution – adding Gaussian noise that is drawn from a standard normal distribution and then scaled by a hyper-parameter  $w$ , to the text vectors during training. Intuitively, this noise helps to close

Soln

<sup>1</sup><https://github.com/allenai/close>

the modality gap by spreading out the text vectors and overlapping them with the image vectors. Figure 3b visually shows that even a small amount of noise leads to much better overlapping of the image and text vector spaces. The noise also encourages the model to be more robust to minor changes or variations to the input vectors, and thus be better prepared for the shift caused by switching from text to image vectors.

how noise helps – 2 points

A second motivation for using random noise is the observation that image vectors capture certain subtle visual details like lighting, background, or camera position that are not reflected in the text vectors. To illustrate this, we show a small case study in Appendix 5 where we observe that semantic changes (e.g., changing the subject of a caption or image from “dog” to “cat”) result in a relatively consistent directional shift for text vectors, but has a more erratic effect on image vectors. Adding noise to the text embedding helps to mitigate this problem by simulating the fact that, even for semantically similar inputs, image and text vectors can still have minor differences due to the additional information encoded in the images.

After adding the noise we re-normalize the vector to unit length to match the image vectors that will be used during evaluation. We study the modality gap and other approaches to handling it in more detail in Section 4.

## 3. Experiments

We report results on four V&L tasks: captioning, visual entailment, VQA and visual news, and when training CLOSE using only text generated by a language model.

### 3.1. Setup

We construct pure-text training datasets for these tasks using the text annotations from the relevant training datasets, and, for some tasks, text captions of the training images. Our primary point of comparison is a CLOSE model trained with the training images, in which case the images are encoded with the image encoder during training in the same manner as done during testing. This model does not experience domain shift, so we view it as an upper bound. We emphasize that in practice the text training data could come from many other possible sources, see Sect. 5 and Sect. 3.3 for additional experiments that demonstrate this, we use these text sources since they closely match the data the models with images are trained on and therefore allow us to better isolate and study what performance is lost due to the image-text domain shift.

We use T5<sub>base</sub> [52] and CLIP<sub>ViT-L/14</sub> [51], a noise level of 0.08, and a fixed set of hyper-parameters for all tasks to demonstrate our method is effective even when there is no image/text validation set to tune on. See Appendix 1 for hyper-parameter details. We additionally show results when the noise level is tuned on validation sets, and

why 4?  
— 4 tasks

observation  
for modality  
gap

reason  
image and  
text pairs  
need to be  
relatively  
closer  
compared to  
other  
unrelated  
images/text  
in the  
semantic  
space. Need  
not be close  
in absolute  
terms

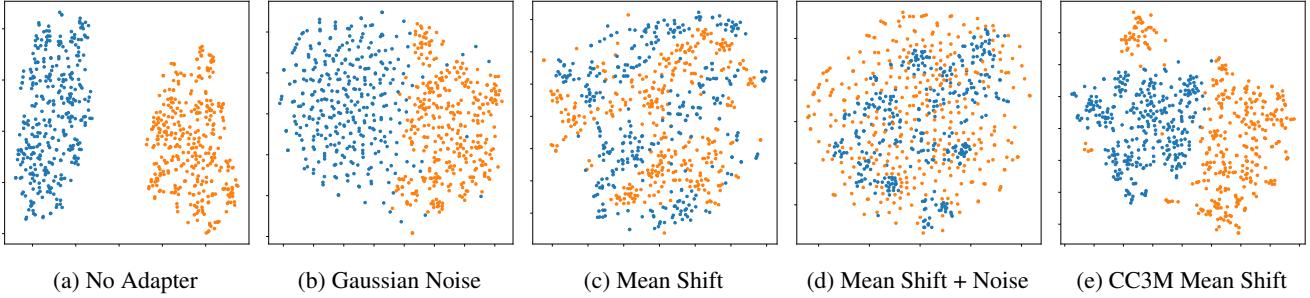


Figure 3: t-SNE [65] plots for various adapters on 350 randomly selected image vectors (blue) and paired caption vectors (orange) from COCO captions. The first two panels demonstrate CLOSE, and the remaining three show additional adapters we study in our analysis (Section 4).

Model	Text-Only	Cap. (Single)	Cap. (Mult.)	VE	VQA	E-VQA	VN
Prior Work	✓	-	ESPER Style [79] 78.2	CLIP Cls. [57] 66.6	TAP-C [57] 38.7	-	-
CLOSE w/o Noise	✓	16.4	68.7	68.2	60.2	59.8	32.1
<b>CLOSE (Ours)</b>	✓	80.5	95.3	75.9	59.6	62.9	80.8
CLOSE w/Tuned Noise		95.4	98.4	75.9	61.9	64.3	80.8
CLOSE w/Images		113.2	113.2	77.7	65.4	67.9	105.7

Table 1: Results on V&L tasks. Models in the last two rows require images and so are upper bounds for CLOSE. We report CIDEr [66] for captioning with single and multiple captions, visual entailment test accuracy, VQA 2.0 test-dev accuracy, E-VQA validation accuracy, visual news test CIDEr. See Appendix 2 for other metrics and more detailed results.

when the noise is removed, to study the effect of noise on CLOSE.

### 3.2. Results

Results are shown in Table 1. Due to space constraints, we only report one metric for each task here and include more results in Appendix 2. We also show the best method from prior work, when present, that does not use images.

**Image Captioning.** For captioning, we use text captions as both the input text and the target output text. However, we find that, if multiple captions about one scene are available, it is beneficial to use different captions about the same image as the input and target text. We call the first setting *captioning (single)* and the second *captioning (multiple)* and evaluate both since they facilitate different training setups. We evaluate on COCO Captioning [6] using the Karpathy split [28]. We train our text-only models using just the captions in the training data. We treat all captions per image as a group for the multiple-caption setting and use each caption individually in the single-caption setting.

CLOSE reaches 95.3 CIDEr in the multiple caption setting, showing high captioning competency despite not using images. In the single caption setting, performance is reduced but can be increased to 95.4 with higher noise levels. Our approach is substantially better than recent zero-shot

methods such as MAGIC (49.3) [61] and Socratic Models (44.5) [81], and is 17 points ahead of ESPER Style (78.2) [79] which also uses text captions.

**Visual Entailment.** Visual entailment requires determining whether a premise image either entails, contradicts, or is neutral with respect to a hypothesis sentence. During training, a text premise is used instead of an image. The hypothesis sentence is always text and is encoded with the language model. We train on SNLI [45] (a language-only dataset) and evaluate on SNLI-VE [74] (a vision and language dataset). Despite not using images, CLOSE achieves similar performance to the image model. Song *et al.* [57] also experiment with this task, but we find adding Gaussian noise allows us to surpass their result by over 9 points.

**VQA.** To train a VQA model we use data that contains a sentence describing a scene (encoded with the text encoder), a question (encoded with the language model), and a target answer. We consider two datasets. First, we pair COCO captions with questions about the same image from VQA 2.0 [17]. However, in this dataset, the questions might ask about details of the image not included in the caption, and thus cannot be answered by the text-only model. Hence we also train and evaluate on VQA-E [34] which contains a subset of the VQA 2.0 questions paired with COCO captions that have been verified to contain the answer.

DOUBT: if  
we have  
multiple  
captions for  
a same  
image then  
is it  
beneficial to  
use that in  
the same  
batch or  
different  
batches?

Model	B-4	M	C	S
MAGIC [58]	12.9	17.4	49.3	11.3
CLOSE w/Coco	29.5	25.6	98.4	18.3
CLOSE w/GPT-J RNG	19.6	20.9	63.2	13.8
CLOSE w/GPT-J Unigram	<b>23.2</b>	<b>22.2</b>	<b>78.9</b>	<b>15.6</b>
CLOSE w/OpenAI Curie	18.5	21.2	69.0	14.9

Table 2: BLEU-4, METEOR, CIDEr, and SPICE on the COCO validation set when training on synthetic captions.

These training sets have significantly different question distributions due to the filtering done in VQA-E, so we evaluate models either on the VQA 2.0 test-dev set or the VQA-E validation set<sup>2</sup> depending on what train set was used. There is no prior work for this task in the text-only setting, however CLOSE does outperform TAP-C<sub>ViT-B/16</sub> [57], a CLIP-based zero-shot approach.

For VQA-E, we observe only a 3.5 point drop in accuracy relative to image training while surpassing the baselines. The gap is more significant on VQA 2.0, which we attribute to the sometimes poor alignment between the captions and questions, although our method is still within 5 points of the model trained on images.

**Visual News.** Visual news requires captioning an image in the context of a news article, and which therefore often requires mentioning the people, locations, and events from the article text [40]. CLOSE is easily extended to this setting by using the caption as both the image text and the target output, while the article is given as additional context to the language model. For this task, we randomly sample 15% of the training data each epoch due to the large dataset size, and use OpenCLIP instead of CLIP since our previous experiments found it slightly improves performance. CLOSE with images achieves over 105 CIDEr, a significant improvement over the previous best benchmark of 50.5 CIDEr [40]. Training without images also outperforms the previous state-of-the-art, obtaining a respectable 80.8 CIDEr. See Appendix 5 for qualitative examples.

**Discussion.** Overall, performance is comparable to the model trained with images showing CLOSE is able to transfer skills between modalities. Tuning the noise level can benefit some tasks, therefore better heuristics for choosing the noise level or leveraging a small image/text validation set could additionally improve performance. On the other hand, removing the noise reduces performance drastically across almost all tasks. This is because the noise plays an important role in addressing the modality gap.

### 3.3. Training with Data from a Language Model

Next, we use CLOSE to train a captioning model on synthetic data generated by a language model. We first con-

<sup>2</sup>VQA-E does not have a test set

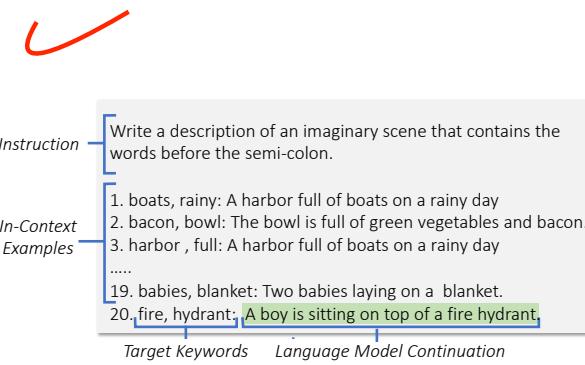


Figure 4: Prompt used to generate a synthetic caption from a language model. The language model’s continuation (highlighted text) is used as a synthetic caption.

struct a prompt that includes a natural language instruction and some example captions following an in-context learning approach [4], shown in Figure 4. To generate a diverse set of captions, we prefix each caption with two keywords that occur in that caption, and end the prompt with two new keywords to be used in the caption to be generated (“fire” and “hydrant” in Figure 4). Then diverse captions can be constructed by changing the ending keyword pair. To reduce the chance of caption style affecting the quantitative evaluation, we take steps to better match the style of the Coco captions, although in settings where the precise style is of less importance this would not be required. We generate 100k examples from three generation methods:

**GPT-J RNG.** Examples are generated using a 6 billion parameter open source language model, GPT-J [68], with 50 in-context examples. Keywords are sampled uniformly at random from keywords in the COCO training data.

**GPT-J Unigram.** Keywords are instead sampled to match the unigram distribution of COCO captions.

**Curie Unigram.** Generations are from OpenAI Curie<sup>3</sup> with 20 examples and unigram-matching.

Results on COCO are shown in Table 2. Our best result achieves 78.9 CIDEr. Inspection shows that, even with our keyword sampling approach, many errors are still caused by style issues, and that style also explains the reduced performance of the Curie model. For example, the synthetic captions from the Curie model are 23 times more likely than the COCO and the GPT-J captions to use the word “opens” (e.g., “a living room that opens onto the balcony”), and use “cellphone” while “cell phone” is much more common in COCO. More details are in Appendix 3. This illustrates how, when using this method, the choice of language model can have subtle effects on the style of captioning that will be learned. Despite this issue, this is still a very strong result that surpasses the zero-shot method MAGIC [58].

<sup>3</sup><https://beta.openai.com/docs/models/gpt-3>

Bias	Mag.	MG	$\Delta$	Cap.	VE	VQA
none	0.0	0.26	1.00	94.4	64.3	75.9
mean	0.8	0.62	0.69	92.8	64.7	75.4
-mean	0.8	-0.10	0.85	84.3	62.0	71.8
RNG	0.2	0.25	0.98	93.5	63.9	75.3
RNG	0.5	0.24	0.89	92.5	64.2	75.3
RNG	0.8	0.20	0.78	89.3	63.7	74.8
RNG	1.0	0.18	0.71	87.2	63.8	74.2
RNG	2.0	0.11	0.45	73.7	61.4	71.3

Table 3: Text vector translation-sensitivity analysis. The first three columns show the translation magnitude, the resulting modality gap on COCO, and the cosine similarity to the original vectors. The following columns show CIDEr captioning score, accuracy on VQA-E, and accuracy on visual entailment on validation sets.

## 4. Analysis

Our approach opens up two intriguing questions: (1) Why does embedding substitution work even when text and image vectors are generally quite far apart? (2) Can methods that leverage additional data to better close the modality gap improve upon this approach? We do two analyses to answer these questions. Furthermore, we study how different choices for the contrastive embedding model or for the language model affect our method’s performance.

### 4.1. Sensitivity Analysis

To help answer the first question, we perform a sensitivity analysis on the input text vectors. To do this, the model is trained while adding a constant vector to the normalized text vectors and then re-normalizing, and tested on the unaltered image vectors as before. This alteration will change how the text vectors are distributed relative to the image vectors, but will not change how the text vectors are distributed relative to one another. We show results when using a random vector (note the same vector is used for all of training, it will just be selected randomly at the start of training) of different magnitudes, the mean difference of text and image vectors to represent a shift towards the image vectors, and the negation of that vector to shift away from the image vectors. In all cases, we continue to add Gaussian noise as before.

Results are shown in Table 3. For random vectors (RNG), we report the average of three runs with 3 different vectors. Overall, we see only minor degradation when using random vectors until very large shifts are used, showing the model is generally insensitive to shifting the text vectors during training. Shifting the vectors towards the images (mean) can result in a slight gain in performance, and shifting the vectors away from them (-mean) results in a more significant decrease, showing the model is not completely insensitive. However it is still notable that vector substitutions work well even as the text vector’s positions

Method	MG	Cap.	VE	VQA	VN
CLOSE	0.26	94.3	75.9	64.3	80.8
+Cov. (Coco)	0.62	<b>106.5</b>	75.5	65.5	<b>84.1</b>
+Cov. (CC3M)	0.58	95.1	75.8	65.0	-
+Linear (Coco)	0.81	99.5	<b>76.0</b>	<b>65.7</b>	-
+Linear (CC3M)	0.75	81.8	75.5	64.9	-

Table 4: Results with adapters built with paired data. The modality gap on COCO captions, captioning CIDEr, visual entailment accuracy, VQA-E accuracy and visual news CIDEr are shown. The last task is more complex and so we only experiment it with one promising adapter.

are significantly randomized.

We hypothesize that this insensitivity is due to two reasons. First, most directions in the shifted feature space are predictive of the output in the same manner as before because the text vectors do not change relative positions. Second, the Gaussian noise trains the model to be insensitive to shifts in unimportant directions in the feature space, which often include the direction of the shift. This insensitivity provides part of the answer to question 1. A major source of the modality gap is a constant shift between the image and text vectors [38]. However, addressing this is not as important as one might expect because CLOSE is not highly sensitive to the absolute positioning of the text vectors.

### 4.2. Learned Adapter Analysis

As suggested by Figure 3c, mean shift might not be perfect at aligning the text and image vectors, so we hypothesize more sophisticated adaption methods could improve performance. More complex adapters generally require a paired image/text corpus to train on, so we avoid using them in our main CLOSE method. However, here we investigate them to better understand how much performance they could potentially contribute. To study the difference between using high-quality annotated data or web data we use both COCO captions and Conceptual Captions 3 Million (CC3M) [54]. For Coco we use the 30k captions from the “restval” set of the Karapathy split, which do not appear in our train, eval or test sets, and for CC3M we use a random sample of 100k image/text pairs. We consider two adapters:

**Linear Adapter.** We learn the modality shift by training a linear model to minimize the Euclidean distance between the adapted text vector and its paired image vector. We continue to add Gaussian noise after applying this model.

**Structured Noise with Covariance Matrix.** Even in principle, we do not expect there to be a perfect one-to-one mapping between text and image vectors because an image vector can be similar to many different texts that describe different parts or details of the image. This motivates us to approach the problem from the perspective of better un-

✓

Egocentric Captions



I saw a bird perched on a sand beach looking at the ocean.  
We walked past a kitchen with a window looking out onto a street.  
We are flying kites in a park.  
My mom is making pancakes in the kitchen.  
We visited an old building with a bicycle leaning against it, next to a brick wall.  
We are playing a video game with controllers in our hands.

Uplifting Captions



A group of people are sitting around a table enjoying pizza and laughter.  
A flock of birds fly overhead as the sun sets in the horizon.  
Two girls sitting on the back of a boat contemplating life's mysteries.  
A bunch of stuff animals on a train journey to reach home.  
A beautiful purple tulip flower pot is in bloom.  
A man skiing down a snowy hill to conquer the high.

Harry Potter Captions



Harry Potter was so excited to start his first year at Hogwarts!  
Gellert Grindelwald looked around at the assembled students and smiled.  
Lucius Malfoy watched with satisfaction as the death eaters gathered around Harry Potter.  
Delores Umbridge sat at her desk, a satisfied smile on her face.  
Rubeus Hagrid roared with laughter as he saw the look of terror on Harry Potter's face.  
Lord Voldemort laughed softly, a cold sound that made the hairs on the back of Harry Potter's neck shiver.

Reviews Captions



A perfect gift for a friend who has a flower garden. The roses are beautiful.  
The leash is well made and easy to put on and take off. My dog is very happy with it.  
This was a wedding cake for my husband and he loved it. He was very happy with the cake.  
This is a great oven. It cooks evenly and is easy to clean. I would recommend it.  
I bought this as a gift for a friend. She loves it. It is very soft and cuddly.  
Fast delivery: I received my order in a timely manner and it was in good condition. I would order from them again.

Figure 5: Examples of stylistic captions produced by CLOSE trained with only text data, and then applied 0-shot to images.

derstanding how text vectors are distributed around its related image vectors, instead of just trying to learn a simple mapping function. In Appendix 4, we provide insight into how the vector differences from COCO image-caption pairs follow a particular shape. To capture this shaped relationship between text and images, we add Gaussian noise whose mean and covariance are learned from the differences between text-image vectors in the auxiliary corpus, to the text during training. This noise is expected to better simulate the text-image shift that will occur during evaluation.

Results are shown in Table 4. We observe large improvements on captioning, modest improvements on VQA and

visual news<sup>4</sup>, and similar performance on visual entailment using the adapters from COCO, with the structured noise approach being significantly better on captioning, and slightly worse on the other tasks. The CC3M adapter also achieves mild gains, although it is less effective. This shows the training data used for the adapter is important, a point that can be qualitatively observed in Figure 3c and Figure 3e.

#### 4.3. Performance Analysis of Different CLIP and T5 Models

Finally, we study how different choices for the contrastive embedding model or for the language model affect

<sup>4</sup>We only test one adapter on this task due to the longer training times

CLIP Model	T5 Model	Cap.	VE	VQA
ViT-L/14	small	94.4	74.9	59.9
ViT-L/14	base	95.4	76.1	64.3
ViT-L/14	large	93.9	75.1	65.2
ViT-B/32	base	91.1	75.3	61.4
RN101	base	90.0	75.4	59.8
RN50	base	90.2	75.3	60.4
RN50×4	base	92.0	75.3	61.5
RN50×16	base	93.4	74.4	62.5
RN50×64	base	96.1	75.8	64.2
OpenCLIP [24]	base	99.2	<b>76.3</b>	65.1
EVA-CLIP [13]	base	<b>101.7</b>	75.53	<b>66.6</b>

Table 5: **Ablations with different contrastive and language models**. The first column indicates which CLIP model was used, with OpenCLIP indicating we use the ViT-L/14 OpenCLIP model trained on Laion 400m [24]. The last three columns show CIDEr on COCO captioning in the single caption setting, accuracy on visual entailment, and overall accuracy on VQA-E on the validation sets.

the performance of our method. Results for captioning, visual entailment, and E-VQA are shown in Table 5. For these experiments we use the tuned noise values in order to compare best-case performance. We find the optimal noise level for these models generally does not change as these components are altered, so we use the same noise levels as our main results for all these experiments.

There is a consistent decrease in performance when using CLIP versions other than ViT-L/14, with only RN50×64 being comparable, showing that CLOSE gains effectiveness as the contrastive model becomes more powerful. We also observe much less dependence on the size of the T5 model, with the large model increasing performance on VQA but not on the other tasks. The OpenCLIP model is generally more effective and boosts the captioning results to nearly 100 CIDEr. The EVA-CLIP model [13] further boosts VQA scores, approaching our main result with images (67.9), showing that CLOSE’s performance can be improved by enhancing the contrastive model.

## 5. Stylistic Captioning

We demonstrate an application of our method by applying it to the task of constructing captions with specific writing styles. Our general approach is to gather text-only training data that exemplifies the style we want the model to use, train on them as if they were text captions as done in Section 3.2, and then apply the model to images. To show that a diverse range of natural language data sources can be used to learn different styles we show four captioning styles, each of which uses a different method of collecting training data.

**Ego-Centric.** Section 3.3 shows that our model can be

trained using data generated by a language model. Now we demonstrate an application of that approach by using the language model to generate captions in an ego-centric style. We use the same prompt format as before (Figure 4), only now with 20 examples of manually authored captions written from a first-person perspective. We again sample keywords randomly from those found in COCO training captions to generate diverse prompts and obtain 20k captions using OpenAI’s GPT-3 model. We apply this model to COCO validation images, shown in the top row of Figure 5, and observe it learns to use a variety of first-person language while accurately describing the image.

**Uplifting.** We use a publicly available dataset [14] to collect 6k examples of uplifting captions (no images). Results are shown in the second row in Figure 5, where we observe the model adds warm and optimistic details to its captions.

**Character-Based.** Next, we target character-based captions that use proper nouns and describe images as if they were from a story. Using proper nouns would be a significant hurdle for many existing systems due to the lack of image/name paired data in existing datasets. However, CLOSE can leverage CLIP’s ability of recognizing names of famous people [51] to handle that problem. We first pick 33 Harry Potter characters. Then only a few excerpts from the Harry Potter books or fan fictions are manually collected and used, together with the characters, as prompts to GPT-3 to create 13k captions. Results on relevant photos are shown in the third row of Figure 5. The model uses the correct names and image content, while sometimes making up plausible events that could give additional context to the image as if it was a scene in a book or a movie.

**Reviews.** We train a model to write captions like a customer writing a review. For training data, we gather publicly-available Amazon product reviews<sup>5</sup> and select positive reviews that are a maximum of 40 tokens long. As shown in Figure 5 bottom row, the captions use a variety of language to write positive reviews of the items in the photos.

## 6. Related Work

**Using Contrastive Models.** Many vision and language contrastive models have been constructed, including CLIP [51], ALIGN [25], UniCL [76] and OpenCLIP [24], and recent multi-modal models that contain a contrastive training component [78, 80, 32]. Typically these models are used either zero-shot, which is effective for image classification but challenging for more complex tasks like captioning or visual entailment [57, 61, 81], or as feature extractors for down-stream tasks [55, 29, 18, 12, 44, 50, 82, 72]. Our work offers a compromise between those two approaches by allowing models to be trained with only textual data, which

<sup>5</sup><https://www.kaggle.com/datasets/bittlingmayer/amazonreviews>

ablation  
shows  
there is  
less  
dependen  
ce on the  
T5(text  
decoder)

substantially improves upon zero-shot performance without requiring annotated images.

**Zero-Shot Vision Using Language Models.** Several recent works have combined large language models with pre-trained vision models to perform vision tasks zero-shot. Methods include using reinforcement learning to learn how to generate text that matches a CLIP Embedding [79], using CLIP to guide inference in the LLM [62], or using a pre-trained model to generate text describing an image to pass into the language model [81]. Compared to these methods our approach of leveraging text training has several advantages. Fine-tuning on text-only data enables our model to learn task-specific details and subtleties that are challenging for fully zero-shot methods, such as the style of captions to be generated. Our approach also works effectively with smaller language models (CLOSE only uses 220M trainable parameters) which significantly reduces the computational demand.

**Cross-Modal Transfer Learning.** Transfer learning has typically focused on transferring skills from one modality to the same modality. CROMA is an exception and uses a modality-invariant feature space to achieve transfer similar to our work, however, it is limited to classification tasks and is few-shot rather than zero-shot [38]. Pre-trained language models have been shown to learn skills that can transfer to new modalities [42], however, this will be ineffective for task-specific skills such as a desired captioning style or learning the space of output labels. Several multi-modal/multi-task models have learned many tasks in different modalities simultaneously [41, 70, 37, 26] and could thus potentially transfer skills between them, with High-MMT in particular showing positive results [37]. Our work studies the more challenging zero-shot setting (meaning no training data in the target modality is available), and therefore requires all the needed skills to be learned from a modality different than the one used in evaluation.

Recently, Song *et al.* [57] use a similar vector-substitution trick with CLIP to train visual entailment models, however they do not use noise or other methods that address the modality gap. Yu *et al.* [79] use reinforcement learning to train a model to generate text that CLIP ranks as being close to input images, and text data to learn captioning styles, although they do not directly train on text versions of the vision tasks. Concurrently with our work, Nukrai *et al.* [48] and Wei *et al.* [35] propose text-only approaches leveraging CLIP with either Gaussian noise similar to CLOSE, or using a projection of the text embeddings. Our work does additional analysis, covers more tasks including experiments using data generated by a language model, and achieves better captioning results.

**Domain Invariant Representations.** Using domain-invariant features to achieve out-of-domain generalization has a long history in transfer learning. Work in this area has

shown such features can be built from multi-domain training data [69, 19], small amounts of labelled data in the target domain [9, 64], and unsupervised data [71, 59]. Methods include using adversarial learning to remove domain-dependent features [16, 36, 63], using maximum mean discrepancy to ensure features are distributed similarly across multiple domains [31, 3] and various data augmentation approaches to prevent models from learning domain-dependent features [85, 84, 67, 53]. The effectiveness of Gaussian noise in making models robust to domain shifts in these features has also been observed in image classification [33]. While we also use domain-invariant features, the domain shift we study is more extreme than what is typically studied due to the change in modalities, and we show large-scale contrastive models can be an effective source of invariant features if used correctly.

**Stylistic Captioning.** Stylistic captioning models can be built by authoring captions of the desired style [46, 14, 21, 56] and applying standard captioning methods. However, since creating such annotations is expensive, many stylistic captioning methods additionally transfer from captions with other styles by pre-training or multi-tasking [46, 47, 77]. Other methods have combined unstylized captioning data with text data in the desired style through methods such as adversarial learning [5], multi-tasking with language modelling [14], or factoring caption writing into style and context components so that the style component can be learned from the text [14, 83]. Most similar to our work, Tan *et al.* [60] train a model to generate text from either images or text using a shared encoding space and learned style embeddings. Unlike these methods, our approach does not require the use of any paired image/caption data.

## 7. Conclusion

We have shown that the multi-modal semantic vector space learned by contrastive models can be used for cross-modal generalization through CLOSE, and studied its sensitivity and what improvements can be made with trained adapters. We have also conducted experiments on multiple vision and language tasks and demonstrated a specific application to stylistic captioning. Beyond stylistic captioning, CLOSE is applicable to many other cases where training data is abundant in one modality but scarce in another. Possible use-cases include: training a captioning model for 3D scenes using image captioning data; training a model to summarize a video using text summarization data; and training a model to perform tasks like VQA or captioning for less-studied modalities like tables, graphs, or sensors without having to annotate additional data for all modalities. As more powerful contrastive models that span more modalities are trained, we expect CLOSE to yield better results and gain more use cases.

## References

- [1] Mohamed Afham, Isuru Dissanayake, Dinithi Dissanayake, Amaya Dharmasiri, Kanchana Thilakarathna, and Ranga Rodriguez. Crosspoint: Self-supervised cross-modal contrastive learning for 3d point cloud understanding. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9902–9912, 2022.
- [2] Peter Anderson, Basura Fernando, Mark Johnson, and Stephen Gould. Spice: Semantic propositional image caption evaluation. In *ECCV*, 2016.
- [3] Karsten M Borgwardt, Arthur Gretton, Malte J Rasch, Hans-Peter Kriegel, Bernhard Schölkopf, and Alex J Smola. Integrating structured biological data by kernel maximum mean discrepancy. *Bioinformatics*, 22(14):e49–e57, 2006.
- [4] Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, T. J. Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeff Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. *ArXiv*, abs/2005.14165, 2020.
- [5] Tseng-Hung Chen, Yuan-Hong Liao, Ching-Yao Chuang, Wan Ting Hsu, Jianlong Fu, and Min Sun. Show, adapt and tell: Adversarial training of cross-domain image captioner. In *2017 IEEE International Conference on Computer Vision (ICCV)*, pages 521–530, 2017.
- [6] Xinlei Chen, Hao Fang, Tsung-Yi Lin, Ramakrishna Vedantam, Saurabh Gupta, Piotr Dollár, and C. Lawrence Zitnick. Microsoft coco captions: Data collection and evaluation server. *ArXiv*, abs/1504.00325, 2015.
- [7] Yen-Chun Chen, Linjie Li, Licheng Yu, Ahmed El Kholy, Faisal Ahmed, Zhe Gan, Yu Cheng, and Jingjing Liu. Uniter: Learning universal image-text representations. *ArXiv*, abs/1909.11740, 2019.
- [8] Jaemin Cho, Jie Lei, Hao Tan, and Mohit Bansal. Unifying vision-and-language tasks via text generation. *ArXiv*, abs/2102.02779, 2021.
- [9] Hal Daumé III. Frustratingly easy domain adaptation. *arXiv preprint arXiv:0907.1815*, 2009.
- [10] Michael Denkowski and Alon Lavie. Meteor universal: Language specific translation evaluation for any target language. In *Proceedings of the ninth workshop on statistical machine translation*, pages 376–380, 2014.
- [11] Benjamin Elizalde, Soham Deshmukh, Mahmoud Al Ismail, and Huaming Wang. Clap: Learning audio concepts from natural language supervision. *arXiv preprint arXiv:2206.04769*, 2022.
- [12] Han Fang, Pengfei Xiong, Luhui Xu, and Yu Chen. Clip2video: Mastering video-text retrieval via image clip. *arXiv preprint arXiv:2106.11097*, 2021.
- [13] Yuxin Fang, Wen Wang, Binhui Xie, Quan Sun, Ledell Wu, Xinggang Wang, Tiejun Huang, Xinlong Wang, and Yue Cao. Eva: Exploring the limits of masked visual representation learning at scale. *arXiv preprint arXiv:2211.07636*, 2022.
- [14] Chuang Gan, Zhe Gan, Xiaodong He, Jianfeng Gao, and Li Deng. Stylenet: Generating attractive visual captions with styles. In *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 955–964, 2017.
- [15] Kavita Ganeshan. Rouge 2.0: Updated and improved measures for evaluation of summarization tasks. *arXiv preprint arXiv:1803.01937*, 2018.
- [16] Yaroslav Ganin, Evgeniya Ustinova, Hana Ajakan, Pascal Germain, Hugo Larochelle, François Laviolette, Mario Marchand, and Victor Lempitsky. Domain-adversarial training of neural networks. *The journal of machine learning research*, 17(1):2096–2030, 2016.
- [17] Yash Goyal, Tejas Khot, Douglas Summers-Stay, Dhruv Batra, and Devi Parikh. Making the V in VQA matter: Elevating the role of image understanding in Visual Question Answering. In *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017.
- [18] Xiuye Gu, Tsung-Yi Lin, Weicheng Kuo, and Yin Cui. Open-vocabulary object detection via vision and language knowledge distillation. *arXiv preprint arXiv:2104.13921*, 2021.
- [19] Ishaan Gulrajani and David Lopez-Paz. In search of lost domain generalization. *arXiv preprint arXiv:2007.01434*, 2020.
- [20] Tanmay Gupta, Amita Kamath, Aniruddha Kembhavi, and Derek Hoiem. Towards general purpose vision systems. *ArXiv*, abs/2104.00743, 2021.
- [21] Danna Gurari, Yinan Zhao, Meng Zhang, and Nilavra Bhattacharya. Captioning images taken by people who are blind. In *ECCV*, 2020.
- [22] Andrey Guzhov, Federico Raue, Jörn Hees, and Andreas Dengel. Audioclip: Extending clip to image, text and audio. In *ICASSP 2022-2022 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 976–980. IEEE, 2022.
- [23] Ari Holtzman, Jan Buys, Li Du, Maxwell Forbes, and Yejin Choi. The curious case of neural text degeneration. *arXiv preprint arXiv:1904.09751*, 2019.
- [24] Gabriel Ilharco, Mitchell Wortsman, Ross Wightman, Cade Gordon, Nicholas Carlini, Rohan Taori, Achal Dave, Vaishaal Shankar, Hongseok Namkoong, John Miller, Hannaneh Hajishirzi, Ali Farhadi, and Ludwig Schmidt. Openclip, July 2021.
- [25] Chao Jia, Yinfei Yang, Ye Xia, Yi-Ting Chen, Zarana Parekh, Hieu Pham, Quoc Le, Yun-Hsuan Sung, Zhen Li, and Tom Duerig. Scaling up visual and vision-language representation learning with noisy text supervision. In *International Conference on Machine Learning*, pages 4904–4916. PMLR, 2021.
- [26] Lukasz Kaiser, Aidan N. Gomez, Noam M. Shazeer, Ashish Vaswani, Niki Parmar, Llion Jones, and Jakob Uszkoreit. One model to learn them all. *ArXiv*, abs/1706.05137, 2017.
- [27] Amita Kamath, Christopher Clark, Tanmay Gupta, Eric Kolve, Derek Hoiem, and Aniruddha Kembhavi. Webly supervised concept expansion for general purpose vision models. In *ECCV*, 2022.

- [28] Andrej Karpathy and Li Fei-Fei. Deep visual-semantic alignments for generating image descriptions. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 3128–3137, 2015.
- [29] Apoorv Khandelwal, Luca Weihs, Roozbeh Mottaghi, and Aniruddha Kembhavi. Simple but effective: Clip embeddings for embodied ai. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 14829–14838, 2022.
- [30] Diederik P Kingma and Jimmy Ba. Adam: A method for stochastic optimization. *arXiv preprint arXiv:1412.6980*, 2014.
- [31] Haoliang Li, Sinno Jialin Pan, Shiqi Wang, and Alex C Kot. Domain generalization with adversarial feature learning. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 5400–5409, 2018.
- [32] Junnan Li, Ramprasaath Selvaraju, Akhilesh Gotmare, Shafiq Joty, Caiming Xiong, and Steven Chu Hong Hoi. Align before fuse: Vision and language representation learning with momentum distillation. *Advances in neural information processing systems*, 34:9694–9705, 2021.
- [33] Pan Li, Da Li, Wei Li, Shaogang Gong, Yanwei Fu, and Timothy M Hospedales. A simple feature augmentation for domain generalization. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 8886–8895, 2021.
- [34] Qing Li, Qingyi Tao, Shafiq Joty, Jianfei Cai, and Jiebo Luo. Vqa-e: Explaining, elaborating, and enhancing your answers for visual questions. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 552–567, 2018.
- [35] Wei Li, Linchao Zhu, Longyin Wen, and Yi Yang. Decap: Decoding clip latents for zero-shot captioning. In *International Conference on Learning Representations*.
- [36] Ya Li, Xinmei Tian, Mingming Gong, Yajing Liu, Tongliang Liu, Kun Zhang, and Dacheng Tao. Deep domain generalization via conditional invariant adversarial networks. In *Proceedings of the European Conference on Computer Vision (ECCV)*, pages 624–639, 2018.
- [37] Paul Pu Liang, Yiwei Lyu, Xiang Fan, Shengtong Mo, Dani Yogatama, Louis-Philippe Morency, and Ruslan Salakhutdinov. Highmmt: Towards modality and task generalization for high-modality representation learning. *ArXiv*, abs/2203.01311, 2022.
- [38] Paul Pu Liang, Peter Wu, Liu Ziyin, Louis-Philippe Morency, and Ruslan Salakhutdinov. Cross-modal generalization: Learning in low resource modalities via meta-alignment. *Proceedings of the 29th ACM International Conference on Multimedia*, 2021.
- [39] Weixin Liang, Yuhui Zhang, Yongchan Kwon, Serena Yeung, and James Y. Zou. Mind the gap: Understanding the modality gap in multi-modal contrastive representation learning. *ArXiv*, abs/2203.02053, 2022.
- [40] Fuxiao Liu, Yinghan Wang, Tianlu Wang, and Vicente Ordonez. Visual news: Benchmark and challenges in news image captioning. *arXiv preprint arXiv:2010.03743*, 2020.
- [41] Jiasen Lu, Christopher Clark, Rowan Zellers, Roozbeh Mottaghi, and Aniruddha Kembhavi. Unified-io: A unified model for vision, language, and multi-modal tasks. *ArXiv*, abs/2206.08916, 2022.
- [42] Kevin Lu, Aditya Grover, P. Abbeel, and Igor Mordatch. Pretrained transformers as universal computation engines. *ArXiv*, abs/2103.05247, 2021.
- [43] Ximing Lu, Sean Welleck, Peter West, Liwei Jiang, Jungo Kasai, Daniel Khashabi, Ronan Le Bras, Lianhui Qin, Youngjae Yu, Rowan Zellers, et al. Neurologic a\* esque decoding: Constrained text generation with lookahead heuristics. *arXiv preprint arXiv:2112.08726*, 2021.
- [44] Huashao Luo, Lei Ji, Ming Zhong, Yang Chen, Wen Lei, Nan Duan, and Tianrui Li. Clip4clip: An empirical study of clip for end to end video clip retrieval and captioning. *Neurocomputing*, 508:293–304, 2022.
- [45] Bill MacCartney and Christopher D. Manning. Modeling semantic containment and exclusion in natural language inference. In *Proceedings of the 22nd International Conference on Computational Linguistics (Coling 2008)*, pages 521–528, Manchester, UK, Aug. 2008. Coling 2008 Organizing Committee.
- [46] A. Mathews, Lexing Xie, and Xuming He. Senticap: Generating image descriptions with sentiments. *ArXiv*, abs/1510.01431, 2016.
- [47] Omid Mohamad Nezami, Mark Dras, Stephen Wan, and Cécile Paris. Senti-attend: Image captioning using sentiment and attention. *ArXiv*, abs/1811.09789, 2018.
- [48] David Nukrai, Ron Mokady, and Amir Globerson. Text-only training for image captioning using noise-injected clip. *arXiv preprint arXiv:2211.00575*, 2022.
- [49] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *ACL*, 2002.
- [50] Jesús Andrés Portillo-Quintero, José Carlos Ortiz-Bayliss, and Hugo Terashima-Marín. A straightforward framework for video retrieval using clip. In *Mexican Conference on Pattern Recognition*, pages 3–12. Springer, 2021.
- [51] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning*, pages 8748–8763. PMLR, 2021.
- [52] Colin Raffel, Noam M. Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. Exploring the limits of transfer learning with a unified text-to-text transformer. *ArXiv*, abs/1910.10683, 2020.
- [53] Shiv Shankar, Vihari Piratla, Soumen Chakrabarti, Siddhartha Chaudhuri, Preethi Jyothi, and Sunita Sarawagi. Generalizing across domains via cross-gradient training. *arXiv preprint arXiv:1804.10745*, 2018.
- [54] Piyush Sharma, Nan Ding, Sebastian Goodman, and Radu Soricut. Conceptual captions: A cleaned, hypernymed, image alt-text dataset for automatic image captioning. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 2556–2565, 2018.

- [55] Sheng Shen, Liunian Harold Li, Hao Tan, Mohit Bansal, Anna Rohrbach, Kai-Wei Chang, Zhewei Yao, and Kurt Keutzer. How much can clip benefit vision-and-language tasks? *arXiv preprint arXiv:2107.06383*, 2021.
- [56] Kurt Shuster, Samuel Humeau, Hexiang Hu, Antoine Bordes, and Jason Weston. Engaging image captioning via personality. *2019 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 12508–12518, 2019.
- [57] Haoyu Song, Li Dong, Wei-Nan Zhang, Ting Liu, and Furu Wei. Clip models are few-shot learners: Empirical studies on vqa and visual entailment. *arXiv preprint arXiv:2203.07190*, 2022.
- [58] Yixuan Su, Tian Lan, Yahui Liu, Fangyu Liu, Dani Yogatama, Yan Wang, Lingpeng Kong, and Nigel Collier. Language models can see: Plugging visual controls in text generation. *arXiv preprint arXiv:2205.02655*, 2022.
- [59] Yu Sun, Eric Tzeng, Trevor Darrell, and Alexei A Efros. Unsupervised domain adaptation through self-supervision. *arXiv preprint arXiv:1909.11825*, 2019.
- [60] Yutong Tan, Zheng Lin, Peng Fu, Mingyu Zheng, Lanrui Wang, Yanan Cao, and Weipin Wang. Detach and attach: Stylized image captioning without paired stylized dataset. *Proceedings of the 30th ACM International Conference on Multimedia*, 2022.
- [61] Yoad Tewel, Yoav Shalev, Idan Schwartz, and Lior Wolf. Zero-shot image-to-text generation for visual-semantic arithmetic. *arXiv preprint arXiv:2111.14447*, 2021.
- [62] Yoad Tewel, Yoav Shalev, Idan Schwartz, and Lior Wolf. Zerocap: Zero-shot image-to-text generation for visual-semantic arithmetic. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 17918–17928, 2022.
- [63] Eric Tzeng, Judy Hoffman, Kate Saenko, and Trevor Darrell. Adversarial discriminative domain adaptation. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 7167–7176, 2017.
- [64] Eric Tzeng, Judy Hoffman, Ning Zhang, Kate Saenko, and Trevor Darrell. Deep domain confusion: Maximizing for domain invariance. *arXiv preprint arXiv:1412.3474*, 2014.
- [65] Laurens Van der Maaten and Geoffrey Hinton. Visualizing data using t-sne. *Journal of machine learning research*, 9(11), 2008.
- [66] Ramakrishna Vedantam, C. Lawrence Zitnick, and Devi Parikh. Cider: Consensus-based image description evaluation. *2015 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 4566–4575, 2015.
- [67] Riccardo Volpi, Hongseok Namkoong, Ozan Sener, John C Duchi, Vittorio Murino, and Silvio Savarese. Generalizing to unseen domains via adversarial data augmentation. *Advances in neural information processing systems*, 31, 2018.
- [68] Ben Wang. Mesh-Transformer-JAX: Model-Parallel Implementation of Transformer Language Model with JAX. <https://github.com/kingoflolz/mesh-transformer-jax>, May 2021.
- [69] Jindong Wang, Cuiling Lan, Chang Liu, Yidong Ouyang, Tao Qin, Wang Lu, Yiqiang Chen, Wenjun Zeng, and Philip Yu. Generalizing to unseen domains: A survey on domain generalization. *IEEE Transactions on Knowledge and Data Engineering*, 2022.
- [70] Peng Wang, An Yang, Rui Men, Junyang Lin, Shuai Bai, Zhikang Li, Jianxin Ma, Chang Zhou, Jingren Zhou, and Hongxia Yang. Unifying architectures, tasks, and modalities through a simple sequence-to-sequence learning framework. In *ICML*, 2022.
- [71] Garrett Wilson and Diane J Cook. A survey of unsupervised deep domain adaptation. *ACM Transactions on Intelligent Systems and Technology (TIST)*, 11(5):1–46, 2020.
- [72] Mitchell Wortsman, Gabriel Ilharco, Jong Wook Kim, Mike Li, Simon Kornblith, Rebecca Roelofs, Raphael Gontijo Lopes, Hannaneh Hajishirzi, Ali Farhadi, Hongseok Namkoong, et al. Robust fine-tuning of zero-shot models. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 7959–7971, 2022.
- [73] Ho-Hsiang Wu, Prem Seetharaman, Kundan Kumar, and Juan Pablo Bello. Wav2clip: Learning robust audio representations from clip. In *ICASSP 2022-2022 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 4563–4567. IEEE, 2022.
- [74] Ning Xie, Farley Lai, Derek Doran, and Asim Kadav. Visual entailment: A novel task for fine-grained image understanding. *arXiv preprint arXiv:1901.06706*, 2019.
- [75] Hu Xu, Gargi Ghosh, Po-Yao Huang, Dmytro Okhonko, Armen Aghajanyan, Florian Metze, Luke Zettlemoyer, and Christoph Feichtenhofer. Videoclip: Contrastive pre-training for zero-shot video-text understanding. *arXiv preprint arXiv:2109.14084*, 2021.
- [76] Jianwei Yang, Chunyuan Li, Pengchuan Zhang, Bin Xiao, Ce Liu, Lu Yuan, and Jianfeng Gao. Unified contrastive learning in image-text-label space. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 19163–19173, 2022.
- [77] Quanzeng You, Hailin Jin, and Jiebo Luo. Image captioning at will: A versatile scheme for effectively injecting sentiments into image descriptions. *ArXiv*, abs/1801.10121, 2018.
- [78] Jiahui Yu, Zirui Wang, Vijay Vasudevan, Legg Yeung, Mojtaba Seyedhosseini, and Yonghui Wu. Coca: Contrastive captioners are image-text foundation models. *arXiv preprint arXiv:2205.01917*, 2022.
- [79] Youngjae Yu, Jiwan Chung, Heeseung Yun, Jack Hessel, JaeSung Park, Ximing Lu, Prithviraj Ammanabrolu, Rowan Zellers, Ronan Le Bras, Gunhee Kim, et al. Multimodal knowledge alignment with reinforcement learning. *arXiv preprint arXiv:2205.12630*, 2022.
- [80] Lu Yuan, Dongdong Chen, Yi-Ling Chen, Noel Codella, Xiyang Dai, Jianfeng Gao, Houdong Hu, Xuedong Huang, Boxin Li, Chunyuan Li, et al. Florence: A new foundation model for computer vision. *arXiv preprint arXiv:2111.11432*, 2021.
- [81] Andy Zeng, Maria Attarian, Brian Ichter, Krzysztof Choromanski, Adrian Wong, Stefan Welker, Federico Tombari, Aveek Purohit, Michael Ryoo, Vikas Sindhwani, Johnny Lee, Vincent Vanhoucke, and Pete Florence. Socratic mod-

- els: Composing zero-shot multimodal reasoning with language. *arXiv*, 2022.
- [82] Xiaohua Zhai, Xiao Wang, Basil Mustafa, Andreas Steiner, Daniel Keysers, Alexander Kolesnikov, and Lucas Beyer. Lit: Zero-shot transfer with locked-image text tuning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 18123–18133, 2022.
- [83] Wentian Zhao, Xinxiao Wu, and Xiaoxun Zhang. Memcap: Memorizing style knowledge for image captioning. *Proceedings of the AAAI Conference on Artificial Intelligence*, 34, 2020.
- [84] Kaiyang Zhou, Yongxin Yang, Timothy Hospedales, and Tao Xiang. Deep domain-adversarial image generation for domain generalisation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2020.
- [85] Kaiyang Zhou, Yongxin Yang, Yu Qiao, and Tao Xiang. Domain generalization with mixstyle. *arXiv preprint arXiv:2104.02008*, 2021.

## Appendix

Model	Mode	B-4	M	C	S
CLOSE w/Images	-	34.4	27.8	113.2	20.4
CLOSE w/Tuned Noise	S	28.6	25.2	95.4	18.1
CLOSE w/Tuned Noise	M	29.5	25.6	98.4	18.3
ESPER Style [79]	-	21.9	21.9	78.2	-
CLOSE w/o Noise	S	4.2	12.2	16.4	6.5
CLOSE w/o Noise	M	21.9	20.6	68.7	13.5
CLOSE	S	22.1	23.7	81.2	17.7
CLOSE	M	29.5	25.7	97.8	18.3

Table 1: Results on the caption test set in single-caption setting and multiple captioning setting, M indicates the multiple caption setting and S indicates the single caption setting.

Model	Yes/No	Num.	Other	All
CLOSE w/Images	83.2	44.8	54.9	65.4
CLOSE w/Tuned Noise	79.4	43.4	51.1	61.9
TAP-CViT-B/16 [57]	71.4	20.9	18.6	38.7
CLOSE	77.1	42.1	48.6	59.6
CLOSE w/o Noise	78.6	40.6	49.0	60.2

Table 2: Results on the VQA 2.0 test-dev set.

Model	Yes/No	Num.	Other	All
CLOSE w/Images	80.4	48.4	64.1	67.9
CLOSE w/Tuned Noise	78.2	46.0	59.5	64.3
CLOSE	74.9	45.2	59.2	62.9
CLOSE w/o Noise	76.8	36.8	53.9	59.8

Table 3: Results on the VQA-E validation set.

## 1. Hyperparameters

For all tasks, we fine-tune our model with the Adam optimizer [30] with a linear decaying learning rate starting at  $3e-4$ ,  $\beta_1 = 0.9$  and  $\beta_2 = 0.999$ , batch size of 128, and train for 8 epochs. We use beam search with a beam size of 5 for evaluations. When tuning the noise level, we select 0.04 for VQA, 0.08 for visual entailment and visual news, 0.14 for captioning in the single caption setting, and 0.04 for captioning in the multiple captioning setting.

## 2. Detailed Results

To facilitate more detailed comparisons with other works, we present results across more metrics of our evaluated datasets. In all tables, upper bounds that use images are shown above the dashed line.

**Captioning.** We present results in Table 1 for BLEU-4 [49],

Model	Val	Test
CLOSE w/Images	77.0	77.7
CLOSE w/Tuned Noise	75.9	75.9
CLIP Classifier [57]	67.2	66.6
CLOSE	75.9	75.9
CLOSE w/o Noise	68.7	68.2

Table 4: Results on the visual entailment test and validation set.

Model	B-4	M	R	C
VNC w/Images [40]	5.3	8.2	17.9	50.5
CLOSE w/Images	9.3	10.9	25	105.7
CLOSE	5.4	8.2	19.7	80.8
CLOSE w/o Noise	2.1	4.9	12.7	32.1

Table 5: Results on the visual news test set.

Model	Individual	Any
OpenAI Curie	58.8	85.0
GPT-J	42.7	81.9

Table 6: How often generated captions contain the target keywords when generating synthetic captions using different language models. The second column shows the success rate for individual generations, and the third column shows how often any caption in the 5 captions generated per a prompt contain both keywords.

METEOR [10], CIDEr [66] and SPICE [2].

**VQA.** We present results by question-type for VQA 2.0 in Table 2 and VQA-E in Table 3.

**Visual Entailment.** We present visual entailment results on the test and dev set in Table 4.

**Visual News.** We present results with BLEU-4 [49], METEOR [10], ROUGE [15] and CIDEr [66] following [40] in Table 5. To the best of our knowledge the previous best reported results is from Liu *et al.* [40] which does not make use of a pre-trained language model like CLOSE does. Qualitative results are show Section 5.

## 3. Generating Synthetic Captions using Language Models

In this section, we give more details about how we generate captions using language models and the results from Section 3.3. When generating captions, we use nucleus sampling [23] at  $p = 0.95$  and a temperate of 1, which we find generally improves results. It is not uncommon for the

Word	Image	Curie Model	COCO Model
pictured (100x)		a sandwich is pictured on a white background. CIDEr: 0.76	a sandwich is sitting on a white plate. CIDEr: 1.29
lays (100x)		a cat lays on a computer keyboard. CIDEr: 0.43	a cat is laying on a laptop computer. CIDEr: 1.94
cityscape (54x)		a clock with a cityscape in the background. CIDEr: 0.44	a clock on the side of a tall building. CIDEr: 1.95
person's (13x)		a tennis racquet is seen in a person's hand. CIDEr: 0.62	a close up of a person with a tennis racket CIDEr: 1.12
sunny (3.5x)		a sunny day with people flying kites. CIDEr: 0.09	a number of people on a beach with a kite CIDEr: 0.98

Figure 1: Examples of words that are over-produced by the captioning model trained on the OpenAI Curie synthetic captions relative to the model trained on the COCO captions. The first column shows the word and how much more common it is across captions generated for images in the COCO validation set. The remaining columns provide an example image and a caption from both models with the CIDEr score computed using human-annotated captions.

caption to fail to contain both input keywords, so we sample 5 captions for each prompt and then select a caption containing the keywords if one exists, and select one randomly otherwise. The in-context example captions are prefixed by randomly chosen words that exist within that caption (excluding stop words), and we use randomly selected captions from COCO training captions as the examples. During sampling, we randomly shuffle both the order of the in-context examples and what keywords are used as prefixes for those examples to improve the diversity of the outputs. If doing unigram sampling, we keep track of the distribution of words found in the captions generated so far, and sample new keywords in proportion to how under-represented they are, while never sampling over-represented words.

Statistics for how often the input keywords are correctly included in the caption are shown in Table 6. The success rate is less than 60%, although selecting from 5 generations brings the success rate up considerably. GPT-J is worse than OpenAI Curie, but sampling extra captions helps make up for this deficiency. Future work could integrate a constrained beam search method to address this difficulty [43].

We find that about 10% of GPT-J captions are not coherent or do not describe a visual scene, while these kinds of

captions almost never occur with OpenAI Curie. Overall, for GPT-J, producing 100k captions took about 50 GPU hours using a NVIDIA RTX A6000. For OpenAI Curie, each generation requires approximately 500 tokens per a query, so the total cost was about 100\$<sup>1</sup>. Both methods are far cheaper than annotating data.

As discussed, we observe stylistic differences occur between models trained on synthetic captions and models trained on COCO captions. A particular issue is that, while unigram sampling prevents words becoming under-represented, it still allows some words to become over-represented if the language model has a natural tendency to generate them. Figure 1 contains some examples where the model trained on OpenAI Curie captions uses words like “pictured”, “lays” or “cityscape” that almost never occur in COCO captions and thus lead to low quantitative scores even when used correctly. Interestingly, we find GPT-J is not as affected by this issue, which likely stems from differences in what data the language model was trained on. Nevertheless, the captions do still correspond well to the image content, as shown by reasonably good captioning scores despite these stylistic issues, showing it is possible to learn captioning using only synthetic data.

## 4. The Relationship Between Image and Text Vectors

We perform a small case study by selecting four image/caption pairs that represent two different semantic changes in terms of animal species and positions (the result is shown in Figure 2) and examine how the image or text vectors shift according to these changes. We observe that text vectors move more consistently when either the species or positions of the animals change. This disparity is likely due to random shifts in image semantics that correlate with conceptual changes in the text, such as subtle alterations in the animals’ appearance, textures, or background.

We further analyze how image and text vectors typically differ by computing the differences between image/text pairs in an auxiliary corpus of COCO. We center these differences and apply PCA. The first two plots in Figure 3 show that the first few PCA dimensions explain a large portion of the variance in these differences, showing that differences often occur in similar directions. We also plot the Pearson correlation coefficient for the most related features in the third plot, showing that a number of these features are highly correlated. Indeed, image/text pairs tend to move in a structured manner that follows a particular “shape”. We capture this subtle relationship by studying the covariance matrix of the differences between text-image vectors. We then modify our Gaussian noise that is added to the text during training to better simulate this co-movement.

<sup>1</sup>At the current rate of 0.002\$ per 1k tokens on 11/16/2022

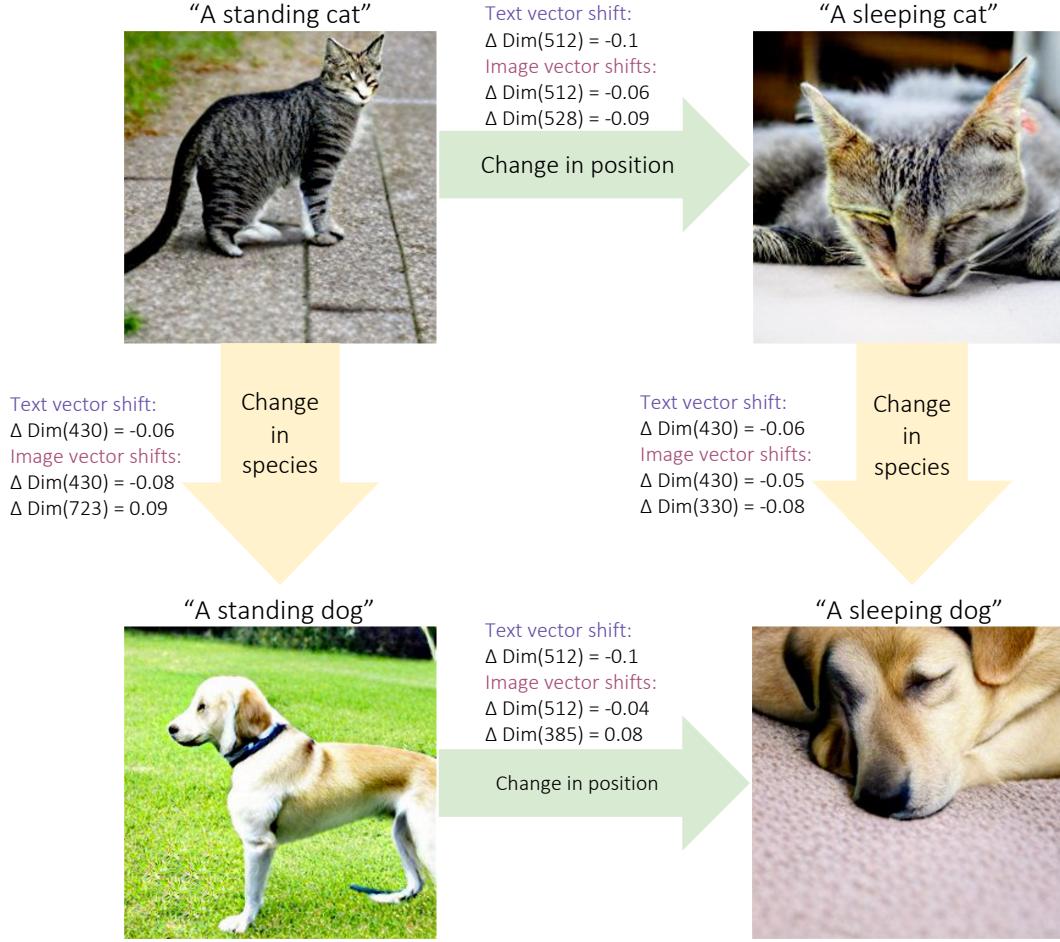


Figure 2: An example of how image/text feature vectors shift with a specific change in species (vertically) or position (horizontally). Text adjacent to each arrow shows any significant changes in the text (purple) or image (red) vector that occurred because of the shift.

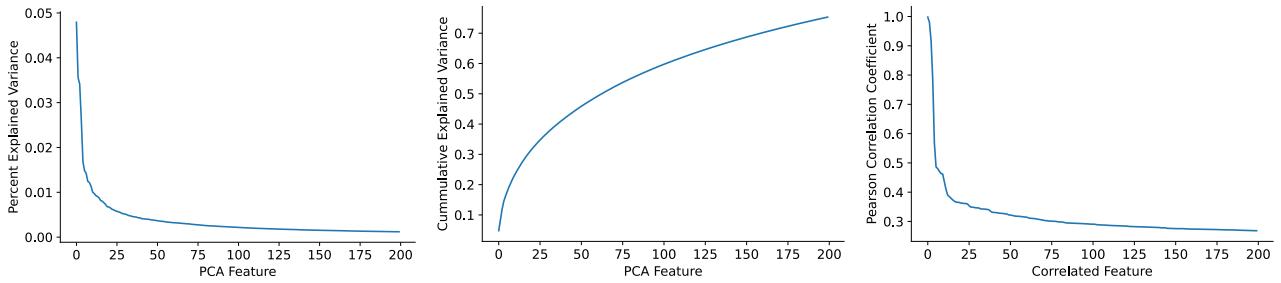


Figure 3: Plots analyzing the differences between image and text vectors for image/caption pairs in COCO captions. Only the first 200 features are shown.

## 5. Visual News Qualitative Examples

We show some qualitative examples for visual news in Figure 4. We observe that close to 50% of time, the predicted captions can be more descriptive (i.e., they can include more details), indicating there is room for this visual

news captioner to grow. There are also some cases in which the predicted captions are better than the ones provided by human (the target captions). But overall, the general sense of both the news images and articles are present in the captions produced by CLOSE.



**target caption:** the trump family cuts the ribbon

**predicted caption:** the trump international hotel opened monday inside the old post office pavilion



**target caption:** lance armstrong waves after receiving the bronze medal in the men's individual time trials at the 2000 summer olympics in sydney

**predicted caption:** lance armstrong was stripped of a bronze medal won at the tour de france in 2000



**target caption:** watch a police officer save a man's life during his lunch break

**predicted caption:** video shows two officers helping a man who reportedly had life-threatening trouble with his foot



**target caption:** national grid has sought to play down the significance of the energy warning

**predicted caption:** national grid has used last resort emergency powers to tell companies to reduce their electricity usage

LocalStaying at the new Trump Hotel in D.C.? You'll pay a price beyond \$700 a night. By Petula DvorakA reservation at the new Trump International Hotel in the nation's capital will carry lots of baggage this fall — and not just the kind you would need to haul around the \$700 a night it is going to cost to stay at the swish new place. Emotionally and politically, the hotel that bears Donald Trump's name and opened Monday inside the Old Post Office Pavilion is already sparking fireworks. And protests right in front of the place. Stay at a Marriott. Book a Hyatt. So what? But consider a reservation at a Trump place — Hotels.com has the D.C. property just down the street from the White House on Pennsylvania Avenue at \$761 for Saturday night — and it gets all kinds of complicated. Endorse the Republican presidential nominee's hotel? "Never. Nope. Not a chance," said Becky Acton, who raised her middle finger at the place as she biked by Sunday night, on the eve of its soft opening. "I would never stay there. No matter what it costs." What about its bar, where wine is sold by the spoon? Or the daily Champagne sabering, where bottles are opened by sword? "No interest," she said. Acton is visiting the District from Columbia, Mo. And she stopped to gape at the Trump hotel as construction workers — many of them Latinos who have been on the receiving end of Trump's slurs against immigrants — rushed around in the dark, tile saws screaming when they cut marble outside the front doors in the final, frantic preparations. She shook her head as she pedaled away. The Klyder family had a different take. "It's beautiful, like a castle," said Emily Klyder, 11, as she photobombed her mom's numerous pictures of the hotel Sunday night.

LONDON — On the day he went public with an admission of doping after years of denials, Olympic officials disclosed one more embarrassment for Lance Armstrong: He was stripped of a bronze medal won at the 2000 Sydney Games. The International Olympic Committee sent a letter to Armstrong on Wednesday night asking him to return the medal, just as it said it planned to do last month. The decision was first reported Thursday by The Associated Press. LANCE: Armstrong's admission part of long-term comeback planOn Monday, Armstrong taped an interview with Oprah Winfrey for broadcast Thursday and Friday on her network. A person familiar with the situation told the AP that the winner of seven straight Tour de France titles confessed to Winfrey to using performance-enhancing drugs. The timing of the IOC move, however, was not related to the TV interview. The IOC executive board discussed revoking the medal in December, but delayed a decision until cycling's governing body notified Armstrong he had been stripped of his seven Tour de France titles and all results since 1998. He then had 21 days to appeal. Now that the deadline has expired, the IOC decided to take the medal away. The letter to Armstrong was also sent to the U.S. Olympic Committee, which would collect the medal. "Having had confirmation from UCI that Armstrong has not appealed the decision to disqualify him from Sydney, we have written to him to ask for the return of the bronze medal," IOC spokesman Mark Adams told the AP. "We have also written to USOC to inform them of the decision." Two months after winning his second Tour de France title in 2000, Armstrong took the bronze in Sydney in the road time trial behind winner and U.S. Postal Service teammate Vyacheslav Ekimov of Russia and Jan Ullrich of Germany.

MLocalWatch police on a lunch break save a man choking on a Subway sandwichBy Justin Wm. MoyerSometimes, a Subway sandwich does not go down smoothly. And when a timely Heimlich maneuver is required, police in Virginia are there to help. That's what Fairfax County police said after posting video of two officers on a lunch break helping a man who reportedly had life-threatening trouble with his foot-long during a June 30 visit to the fast-food chain. "Officer Mulhern and Officer Weaver were taking a quick lunch break in a local Subway," the statement, posted to YouTube, said. "A man approached them who was in distress. Officers noticed that the man was choking and sprung into action." As the video's title put it, the officers were "trained and always ready." "Officer Mulhern utilized his training and administered back blows and the Heimlich maneuver," the statement said. "Due to the officers' quick actions, the man's airway was cleared and all were able to finish enjoying their \$5 footlongs." The man, who was not identified, survived — and got a friendly pat on the back from Officer Mulhern.

National Grid has for the first time used "last resort" emergency powers to tell companies to reduce their electricity usage in an effort to avoid the risk of blackouts. It asked firms to reduce their power demand immediately, issuing a so-called demand-side balancing reserve notice to companies that have signed a contract to say they will take part in the demand reduction scheme. A spokesman said this measure had never been used before, while the grid has previously said it would "only be used as a last resort, after all other actions available in the market have been exhausted". Earlier on Wednesday, National Grid issued an urgent request for energy companies to make more power available after multiple breakdowns at UK power stations. Power firms were asked to supply an extra 500 megawatts between 4.30pm and 6pm, a period when power demand surges, with some people still at work and others arriving home and turning the lights on. The owner of Severn power station, Calon Energy, sold electricity to the National Grid at £2,500 per megawatt hour during the afternoon, industry sources confirmed, compared with the typical price at that time of about £60. National Grid issued the original request by sending a "notification of inadequate system margin", a warning that there was not enough power in reserve to keep the lights on in the event of an unforeseen emergency. Shortly before 6pm, National Grid issued a further statement saying suppliers had responded to its urgent request and 40MW of extra power had been ordered, so the NISM had been withdrawn. "This is one of the routine tools that we use to indicate to the market that we would like more generation to come forward for the evening peak demand period," the company said. "The issuing of a NISM does not mean we were at risk of blackouts. It means that we needed the safety cushion of power in reserve to be higher."



**target caption:** this dec 16 photo shows president obama pausing during a speech at an interfaith vigil for the victims of the sandy hook elementary school shooting in newtown conn

**predicted caption:** president obama must decide where to go big on gun control



**target caption:** tourists follow the pathway

**predicted caption:** a view of the san cristobal bridge in the sierra nevada

WASHINGTON — It's hardly a secret that Barack Obama, like every president no doubt, muses about his ultimate legacy and spot in the presidential pantheon. He approaches his second term confronting tough and shifting challenges that will play big roles in shaping the rest of his presidency and his eventual place in history. In the coming months, Obama will have to decide where to be ambitious, where to be cautious, and where to buy time. He draws political strength from his surprisingly easy re-election in a bad economy. It's partly offset, however, by Republicans' continued control of the House, plus their filibuster powers in the Senate. Some of the big issues awaiting the president's decisions are familiar, long-simmering problems. They include immigration and the need for a tenable balance between taxes, spending and borrowing. Another issue, gun control, jumped to the national agenda's top tier this month following the massacre of first-graders and teachers in a Connecticut school. And the issue of climate change remains unresolved. Veteran politicians and presidential historians say it's almost impossible for Obama to "go big" on all these issues. Indeed, it might prove difficult to go big on even one. While some counsel caution, others urge the president to be as bold and ambitious as possible. "Americans are yearning for leadership," said Gil Troy, a presidential scholar at McGill University. As a president dealing with policy, he said, Obama has generally failed to give "that visionary, powerful address that we came to know and love and expect in the 2008 campaign." Rather than let Congress take the lead on big issues, as it did in drafting the 2009 health care overhaul, Obama should be more forceful in pushing new legislation or using his executive powers.



**target caption:** very few developers have spent time making apps for google glass

**predicted caption:** google's glass app allows users to track their progress in real time

In California, the drought is so much bigger than not being able to water your lawn. We've heard about California's historic drought for years, but today the game changed. While standing on a patch of dry grass in the Sierra Nevada that should have been a snowpack, California Gov. Jerry Brown announced the state's first-ever mandatory cuts in water usage. The state has been working to trim water use since Brown proclaimed a drought emergency last year, but it wasn't enough. More than 98% of the state remains in some level of drought. The water restrictions will affect everything from golf courses to public streets. Campuses, cemeteries and other large landscapes are going to have to make significant cuts in water use. Fifty million square feet of lawns throughout the state will have to be replaced with drought-tolerant landscaping. Families in homes where wells have run dry will have to be relocated. "It's a different world," Brown said. Welcome to California's new normal. What's in #TheShortList: The "religious freedom" bill in Arkansas has divided the governor's family. How much top NCAA basketball coaches get paid • Controversy around video purported to be from inside doomed Germanwings flight • What it really means to be "smartphone-dependent" Short on time? Listen to the audio version of #TheShortList: Arkansas governor's son urged him not to sign 'religious freedom' bill. The "religious freedom" bill in Arkansas is so divisive, it's even split Gov. Asa Hutchinson's own family: His son Seth joined the state's growing opposition to the bill and signed the petition urging him to veto it. Today, the governor said he won't sign the bill in its current form. "It has been my intention all along to have House Bill 1228 to mirror the federal act," Hutchinson said. "The bill that is on my desk at the present time does not ... mirror the federal law." He was referring to the federal Religious Freedom Restoration Act signed by then-president Bill Clinton in 1993.

More than two years after it first introduced Glass to the world, Google is bringing the futuristic device to the UK. But although software developers have had many months to cook up ideas for the eyewear, there are still just a few dozen apps available. Some analysts say the controversial spectacles lack a "killer app" - the one function that will make the average user rush out and buy a pair. Ben Wood, from CCS Insight, says he sees Glass as a "science project," and a "window into what's possible in the future," but by no means a commercial product. But some developers have been giving it a good go, using Glass to... Help the hard of hearing Students at Georgia Tech university have created an app that recognises speech and turns it into on-screen captions in real time. Conversations appear in text on Glass, allowing the wearer to read what is being said, if they didn't quite catch it the first time. The same boffins are working on a similar app that will be able to translate languages in real time. Fantastique. Put out fires As Patrick Jackson, a firefighter from North Carolina in the US, exhibited in a video, emergency service personnel using Glass will be able to summon up critical data such as floor plans and aerial imagery before they enter a burning building. A company called Mutualink is also testing an app that will allow medics to view a patient's medical records as they arrive on the scene of a accident, and police could use Glass to view footage from security cameras, or record altercations. Make your run more fun Race Yourself, an augmented reality app, aims to make running more fun by letting you compete against a version of yourself. A 3D avatar running at the speed of your last run appears in Glass, allowing the wearer to gauge their progress in real time. You can even race against friends and celebrities, or, if you dare, against rampant zombies. Enhance your gallery visits Knowing very little about art need not be a barrier to enjoying a cultural trip.



**target caption:** responding to ad blocking google

**predicted caption:** google's accelerated mobile pages project is at its early stages

Google is attempting to counter the threat from ad-blocking and rivals Facebook and Apple by radically improving the loading speed of web pages on smartphones and tablets. Accelerated Mobile Pages aims to simplify the structure of mobile web pages and place the data needed to deliver them closer to users both physically and virtually in a bid to achieve almost "instant" delivery of articles to anywhere in the world. The project is at its early stages and while the company says it hopes to launch AMP next year, no date has been set. Google said it was unveiling the project early because it wants to work with publishers, the advertising industry and other web platforms such as Facebook and Twitter. The code for AMP will be made public, allowing any company or organisation to tailor it to their own needs. The company said it wanted to collaborate with the wider industry to develop a framework for what works effectively on mobile devices. Early demos of AMP articles show pages with less clutter, such as related stories, but that could change as the project progresses. Google's head of news and social products Richard Gingras said: "This is about making the world wide web great again ... to make sure all users can access the vibrant ecosystem of the world wide web, and get it in a near instant fashion everywhere in the world." The initial plan for AMP has come from the Digital News Initiative, a collaboration between Google and eight European publishers including the Guardian, Les Echos in France, El País in Spain and the Daily Mail in the US. The shift away from desktops to mobile devices has left many publishers struggling to get their articles and other content on to people's devices quickly. Those slow speeds are blamed for the rise of ad-blocking software, which stops many publishers from earning money from advertising online.

Figure 4: Examples of visual news captions produced by CLOSE trained on text captions and news articles alone, and then applied zero-shot to news images and articles.