

Automated Profile Inference with Language Model Agents

Yuntao Du
Purdue University
West Lafayette, USA
ytd@purdue.edu

Zitao Li
Alibaba Group
Seattle, USA
zitao.l@alibaba-inc.com

Bolin Ding
Alibaba Group
Seattle, USA
bolin.ding@alibaba-inc.com

Yaliang Li
Alibaba Group
Seattle, USA
yaliang.li@alibaba-inc.com

Hanshen Xiao
Purdue University
West Lafayette, USA
hsxiao@purdue.edu

Jingren Zhou
Alibaba Group
Seattle, USA
jingren.zhou@alibaba-inc.com

Ninghui Li
Purdue University
West Lafayette, USA
ninghui@purdue.edu

Abstract

Impressive progress has been made in automated problem-solving by the collaboration of large language models (LLMs) based agents. However, these automated capabilities also open avenues for malicious applications. In this paper, we study a new threat that LLMs pose to online pseudonymity, called automated profile inference, where an adversary can instruct LLMs to automatically scrape and extract sensitive personal attributes from publicly visible user activities on pseudonymous platforms. We also introduce an automated profiling framework called AutoProfiler to assess the feasibility of such threats in real-world scenarios. AutoProfiler consists of four specialized LLM agents, who work collaboratively to collect and process user online activities and generate a profile with extracted personal information. Experimental results on two real-world datasets and one synthetic dataset demonstrate that AutoProfiler is highly effective and efficient, and can be easily deployed on a web scale. We demonstrate that the inferred attributes are both sensitive and identifiable, posing significant risks of privacy breaches, such as de-anonymization and sensitive information leakage. Additionally, we explore mitigation strategies from different perspectives and advocate for increased public awareness of this emerging privacy threat to online pseudonymity.

CCS Concepts

• Security and privacy;

Keywords

privacy attack; anonymization; profiling

ACM Reference Format:

Yuntao Du, Zitao Li, Bolin Ding, Yaliang Li, Hanshen Xiao, Jingren Zhou, and Ninghui Li. 2018. Automated Profile Inference with Language Model Agents. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 18 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

In recent years, large language models (LLMs) have become increasingly capable, and autonomous agents utilizing LLMs offer promising opportunities to enhance and replicate complex human workflows. These LLM agents have demonstrated capabilities in diverse areas such as software engineering [55], human behavior simulation [41], and even assisting scientific discovery [5]. However, these same capabilities have raised concerns due to the potential for malicious applications, including social engineering [32] and website exploitation [17]. Notably, there has been a rise in privacy concerns regarding LLMs. In addition to widely studied data privacy risks, such as training data memorization [7, 35], LLMs can also violate individuals' privacy in unexpected ways. For instance, recent studies [37, 47] show that an adversary can use LLMs to steal or leak users' private information by steering chat conversations.

In this paper, we study a new privacy threat that LLMs pose to online pseudonymity. As shown in Figure 1, an adversary can, with the help of LLMs, automatically extract sensitive personal information from the publicly visible online account activities (e.g., posts and comments) of a user on a pseudonymous platform (e.g., Reddit). When a user has conducted substantial online activities, the adversary can even build a detailed description of the user. We call this attack **automated profile inference**. Unlike previous profiling approaches [13, 16] that require significant manual effort and expert knowledge, automated profile inference relies solely on publicly available online activities of the pseudonymous user. The adversary does not need any background information about the user or expertise in profiling. Despite the weaker assumptions about the adversary, we find that the resulting privacy-infringing inferences can reveal highly private information about the user, and the inferred profiles can be exploited to facilitate privacy breaches, such as de-anonymization.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
Conference acronym 'XX, Woodstock, NY

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.
ACM ISBN 978-1-4503-XXXX-X/2018/06
<https://doi.org/XXXXXXX.XXXXXXX>




Figure 1: Illustration of automated profile inference. An adversary instructs AutoProfiler to autonomously scrape¹ and analyze target users’ online activities, extract personal attributes from these activities, and generate detailed user profiles that may cause privacy breaches.

Since users’ online activities on the pseudonymous platform are usually ambiguous, inconsistent, and full of superficially insensitive information, we find that simply feeding these texts to an LLM and instructing it to generate a profile, as explored in [47], is struggling to extract implicit personal details. To address this, we propose an LLM-based multi-agent profiling system, AutoProfiler, which automatically collects and processes noisy online activities, and generates a profile with extracted personal information. Inspired by the well-established methodologies in offender profiling [13], AutoProfiler breaks down the profiling task into four specific components, each managed by an LLM agent with diverse skills and expertise: (i) *Strategist*, who coordinates the overall process and gives instructions to other agents; (ii) *Retriever*, who collects user activities along with relevant context; (iii) *Extractor*, who examines user activities to extract personal details; and (iv) *Summarizer*, who evaluates and refines inferred data to resolve inconsistencies and enhance reliability. By organizing the agents within an iterative workflow, AutoProfiler can autonomously scrape online data, analyze, and extract user profiles without any human involvement. **Effectiveness and Efficiency.** We evaluate AutoProfiler on two pseudonymous platforms (*i.e.*, Reddit and Twitter) to verify the practical feasibility of automated profile inference. We find that AutoProfiler can effectively extract a substantial range of personal information from online activities, covering identifiable attributes such as gender and occupation to more sensitive details, including health conditions and relationships. Notably, AutoProfiler achieves this at an unprecedented scale and cost-efficiency: with over 120× faster processing time and 50× less financial cost than human profilers. We also benchmark AutoProfiler on a human-curated synthetic dataset [56], where it achieves near-expert human performance, with over 87% average accuracy in attributes prediction, significantly surpassing the state-of-the-art LLM-based approach.

Emerging Threats. The effectiveness and efficiency of AutoProfiler in retrieving and distilling personal information make automated profile inference not only feasible but also can be executed on a web scale. This capability presents a new threat to online pseudonymity, as adversaries could exploit inferred profiles to cause severe privacy breaches, such as de-anonymization. Furthermore, the inferred profiles contain deeply private details, which may lead to serious cybercrimes like doxing and cyberbullying [12].

- **De-anonymization.** We demonstrate that the inferred profiles contain substantial personally identifiable information (PII) and

may cause serious privacy breaches, such as de-anonymization. For instance, by linking inferred attributes (*e.g.*, education and occupation) to public profiles on LinkedIn, it is possible to identify the real identities of some pseudonymous Reddit users.

- **Sensitive Personal Information Leakage.** We find that AutoProfiler can capture subtle clues and implications from online activities and uncover deeply private information beyond PII, which we term Sensitive Personal Information (SPI). For instance, an individual may share their experience of struggling with depression and express sympathy to others on forums. These seemingly minor details can be pieced together to reveal the victims’ mental health history, which they may not intend to disclose publicly. Through a quantitative analysis of inferred profiles from 250 Reddit users, we observe that AutoProfiler can expose significant amounts of SPIs, which could be exploited to pose even more severe privacy risks.

We notice that recent work [47] also leverages LLMs to infer personal information. However, it is specifically designed for PII extraction, treating this task as a classification problem with predefined PII categories within synthetic text data. In contrast, our approach addresses a more practical scenario, where LLM agents autonomously scrape, analyze, and infer potential personal information from real-world user activities beyond predefined PII, revealing higher risks of de-anonymization and SPI leakage. We provide a detailed discussion and comparison with this work in Section 5.4.

Potential Mitigations. Given the severity of this threat, we also explore and discuss potential mitigation strategies from different directions. From the user side, public awareness of these emerging threats should be raised, and developing tools to help users better understand and manage their online presence could be valuable. From the platform side, we recommend stronger pseudonymity controls, such as features that allow users to manage the visibility of their activities and adopt different pseudonyms to obfuscate their online personas. Furthermore, LLM providers may incorporate new alignment strategies to detect and prevent such malicious use, and new privacy regulations might be needed to restrict the abuse of LLMs. Finally, we advocate for the privacy research community to develop new privacy-enhancing technologies to address this threat.

Main Contributions. Our key contributions are as follows:

- We introduce automated profiling inference, a new privacy threat to online pseudonymity.
- We propose AutoProfiler, an LLM-based profiling framework which enables autonomously scraping, analyzing, and building user profiles from online activities using specialized LLM agents.
- We evaluate the effectiveness of AutoProfiler on two pseudonymous platforms (*i.e.*, Reddit and Twitter) and showcase that the inferred profiles could cause severe privacy breaches such as the de-anonymization of real Reddit users.
- We conduct a comprehensive evaluation of AutoProfiler using five popular LLMs. Experimental results show that AutoProfiler can achieve high accuracy and low cost and outperform the state-of-the-art LLM-based method.

Responsible Disclosure. Prior to publishing this work, we disclosed our results to major LLM providers. We have also notified Reddit/X about the potential de-anonymization risks of users who use their platform. We have discussed our work with the Institutional Review Board (IRB) and received approval. We refer to Section 7 for a further discussion of ethical considerations.

¹In our experiments, we use the official APIs of platforms to obtain users’ online activities instead of scraping to ensure compliance with the platforms’ terms of service.

2 Background & Related Work

Online Pseudonymity. Online pseudonymity, where individuals interact using pseudonyms rather than their real identities, is a unique characteristic of modern internet culture [24]. Many pseudonymous platforms, such as Reddit and Twitter, allow users to engage under fictitious names. Online anonymity has long been regarded as a fundamental factor in protecting private information and reducing the inherent risks of the web [45], and is widely advocated by both media [46] and research communities [24]. A famous example is the cartoon published in The New Yorker [48], which proclaimed “On the Internet, nobody knows you’re a dog.”

Attribute Inference Attack. The goal of an attribute inference attack is to infer sensitive attributes of target users or records using auxiliary information. Prior studies [20, 30, 57] have shown that online behaviors, such as Facebook likes, can be exploited to infer sensitive attributes (e.g., gender and political views) on social networks. Some research [28, 36] has demonstrated that machine learning models may inadvertently reveal sensitive or proprietary information about their training data. More recently, studies [47] have explored using LLMs for PII extraction. While this line of work focuses on predicting a few predetermined attributes, our approach aims to build a comprehensive profile that potentially includes a broad range of detailed personal information. As a result, these studies are often addressed as classification problems, whereas we approach ours as an inference task.

Profiling. Profiling is the process of constructing a picture of an individual by gathering information about their characteristics, behaviors, patterns, and tendencies. There are various types of profiling, each tailored to a specific purpose. For instance, author profiling [16, 43] aims to identify specific attributes of an author through analysis of written texts, while criminal profiling [6] is a legal tool employed by law enforcement to identify criminals by examining behavioral and psychological traits. In the context of privacy, GDPR [44] defines profiling as the use of personal data to evaluate certain aspects of a natural person. Although these profiling approaches share similarities with ours, our work specifically focuses on automatically inferring the personal implications of publicly available online activities.

LLM Inference and LLM Agent. With the scaling of model and data sizes, LLMs demonstrate impressive inference abilities through in-context learning [11, 40], enabling them to quickly adapt to new tasks through prompting, eliminating the need for fine-tuning process. Building on this capability, LLM-based autonomous agents have garnered significant interest in both industry and academia [52]. Many works have improved the problem-solving abilities of LLMs by integrating discussions among multiple agents, such as code generation [25, 55], human behavior simulation [41] and scientific discovery [5].

LLM for Malicious Use. Beyond inherent vulnerabilities of LLMs like training data memorization [7], recent studies [51] have shown that their inference capabilities pose new threats to security and privacy. Some research [37] demonstrates that LLMs can understand the nuanced implications of conversations and leak personal information during user interactions. Additionally, the autonomous abilities of LLM agents further enable them to engage with real

or virtual environments, facilitating cyberattacks such as website exploitation [17] and social engineering [32].

3 New Threat: Automated Profile Inference

In this section, we introduce the data used (*i.e.*, pseudonymous activities) for automated profile inference and formulate the threat model, followed by outlining its implications for privacy breaches.

3.1 Online Pseudonymous Activities

Online pseudonymity conceals users’ real identities, which fosters an environment where people feel more comfortable expressing thoughts and sharing personal experiences [31]. It has become increasingly common for people to share life experiences, discuss personal issues, and seek advice in online pseudonymous platforms (e.g., Reddit and Twitter/X), resulting in abundant digital footprints. In this paper, we focus on *textual* activities (*i.e.*, posts and comments) on pseudonymous platforms, as these represent the most common and easily accessible data for the adversary.

3.2 Threat Model

We assume that an adversary has access to the online activities (*i.e.*, posts and comments) of a target pseudonymous user u . The adversary’s objective is to construct a detailed user profile D_u based on these activities. Specifically, we make the following assumptions regarding the pseudonymous user and the adversary:

- **Visibility of Online Activities.** We assume that a user’s online activities are visible to the adversary. Currently, this holds even when the platform is not actively helping the adversary. For instance, Reddit does not allow users to hide their activity history; all interactions remain visible unless the user explicitly deletes them. Similarly, on social media platforms like Twitter, an adversary can easily view the posts of any public account simply by following them.
- **Random Usernames.** The user’s username is assumed to be random and unrelated to their real identity. Users may employ different usernames across various platforms to enhance their privacy.
- **Use of off-the-shelf LLMs.** We assume the availability of ready-to-use large language models (LLMs), either through commercial APIs (e.g., OpenAI) or by deploying locally pre-trained models, such as Llama-3 [14]. In Section 5.3, we show that these LLMs have become exceptionally affordable for conducting such profiling attacks.

It is worth mentioning that the adversary does not require expertise in profiling or a deep understanding of specific topics the target user interacts with. All profiling tasks will be automated and performed by LLMs (will explain in detail in the next section).

Profiling Objectives. The abundance of online activities provides substantial behavioral data, which adversaries can exploit to infer users’ traits, a process known as *profiling*. The targeted information for profiling is broadly categorized into two types:

- **Personally Identifiable Information (PII).** This category includes attributes that can be directly linked to an individual, such as {type: “Gender”; value: “Male”}. PII is a well-researched area in privacy studies [35] and characterized by privacy frameworks, such as GDPR [44], HIPAA [2], and CCPA [39].
- **Sensitive Personal Information (SPI).** These attributes are highly sensitive but may be hard to identify individuals, such as {type: “Mental health”; value: “The user has a childhood trauma...”}. The

pseudonymous nature of online platforms encourages users to discuss personal narratives, leading to a significant amount of SPIs. **Breaching Privacy.** In this work, we focus on two potential privacy threats posed by profiling to online pseudonymity: de-anonymization and sensitive personal information leakage.

- *De-anonymization.* Some literature [19] demonstrated that at least 60% of the U.S. population could be uniquely identified using only a few pieces of PII (a subset of the output in the profile inference), such as gender, zip code, and date of birth. However, the inferred profile from a user with pseudonymous activities brings *additional* privacy risks: when more unrestricted attributes (PII or SPI) are revealed, an attacker can use these to cross-reference with auxiliary datasets, further narrowing down the range of the user's real identity. In addition, unlike previous de-anonymization attacks [22, 38, 49] that rely primarily on improperly released private data, profiling from online pseudonymous activities presents a more proactive and powerful attack: we use a case study (Section 5.2) to demonstrate that this threat is not only feasible but also highly automated and scalable with only *public information* and AutoProfiler.

- *Sensitive Personal Information Leakage.* Inferred profiles can contain a wide range of sensitive information (SPI), such as health conditions, relationships, and personal secrets, as shown in Section 5.2. Although this information may not directly identify an individual, attackers can still exploit it to humiliate, intimidate, or threaten the target user, potentially leading to serious cybercrimes, such as cyberbullying (detailed discussion in Appendix E.1).

4 A Framework of Automated Profile Inference

In this section, we introduce AutoProfiler, an automated profiling inference framework, by outlining the challenges and illustrating the details designs and functionalities in AutoProfiler.

4.1 Challenges & Considerations

Noisy Information in Users Activities. Leveraging activities like posts and comments from real-world online interactions for profiling presents several challenges:

- *Irrelevance.* A significant proportion of user-generated content is unrelated to the user's identity. Users engage in a wide range of topics, and much of their activity reveals little to no personal information. This requires filtering out irrelevant content and isolating personal information to construct an accurate profile.
- *Obscurity.* Users often avoid disclosing explicit personal details to protect their identity on pseudonymous platforms. Additionally, interactions are typically informal, necessitating an understanding of contextual nuances in conversations. This leads to indirect and ambiguous clues, which are challenging to extract.
- *Inconsistency.* The behavior of pseudonymous users can be inconsistent or even contradictory, a phenomenon recognized by psychologists as the online disinhibition effect [31]. In pseudonymous environments, individuals may feel less accountable and thus present varied versions of themselves across contexts. For instance, a user might discuss living in Seattle as if they are a local, despite never having resided there. Such inconsistencies make it difficult for LLMs to create a coherent and reliable profile.

Deficiencies of Simple LLM Calls. Given the inherent noise of online activities, we find that simply feeding these texts to an LLM and instructing it to produce a profile is ineffective, as demonstrated

in Section 5.3. Furthermore, users' activities are often extensive and may exceed the context window limitations of LLMs, leading to truncated data and incomplete attribute inference. In addition, automated profile inference involves multiple processes, including scraping, analysis, and inference, which makes it challenging for a single LLM to handle all these tasks effectively.

Unclear Instructions and Demonstrations for LLMs. Many studies have shown that LLMs can quickly adapt to downstream tasks via *prompts* without model finetuning. This adaptability is recognized as one of the emerging capabilities of LLMs [53]. Typically, an effective prompt includes a task-specific description and a few textual demonstrations to guide LLMs in performing a task [11].

However, providing handcrafted examples is challenging due to the complexity of our tasks. For instance, since we do not know what users might discuss online or what personal information could be shared in real-world interactions, it's difficult to create suitable demonstrations for LLM inference. Additionally, even when following popular LLM agent approaches [41, 55] and breaking complex profiling tasks into smaller sub-tasks and assigning them to specialized LLM agents, it's still hard to anticipate all possible scenarios. This makes it difficult to provide clear instructions and examples to help the agents cooperate and effectively use the results from one another.

4.2 AutoProfiler

We propose an LLM-based multi-agent profiling framework, AutoProfiler, to address the above challenges. Specifically, 1) we follow the key processes of offender profiling [13] and decompose automated profile inference into smaller, specific tasks, each managed by specialized LLM agents with diverse skills and expertise. 2) We design an iterative workflow that enables agents to scrape, analyze, and infer from users' activities sequentially. 3) In addition, we devised structured protocols and memory mechanisms to facilitate agents' communication and prevent information overload. These strategies empower agents to collaborate effectively, constructing detailed user profiles autonomously by scraping and analyzing users' online activities.

Roles of Agents. Offender profiling is an investigative strategy used by law enforcement agencies to identify likely suspects by analyzing their behavior and characteristics, which shares many similarities with the goals of our task. The process of offender profiling generally comprises four key stages: decision-process models, background input, assessment, and profile construction [13]. Drawing inspiration from this framework, we define four corresponding roles for agents in AutoProfiler: Strategist, Extractor, Retriever, and Summarizer. The responsibilities of each role are as follows:

- *Strategist* coordinates the attack plan and gives instructions to other agents based on the available information and attack progress.
- *Retriever* gathers the user's activities through publicly available APIs provided by platforms to support Extractor's analysis.
- *Extractor* conducts an in-depth analysis of the user's activities, and extract personal attributes from data collected by Retriever.
- *Summarizer* addresses inconsistencies, contradictions, and duplications in the inferred attributes, refining the results to generate a more reliable profile of the anonymous user.

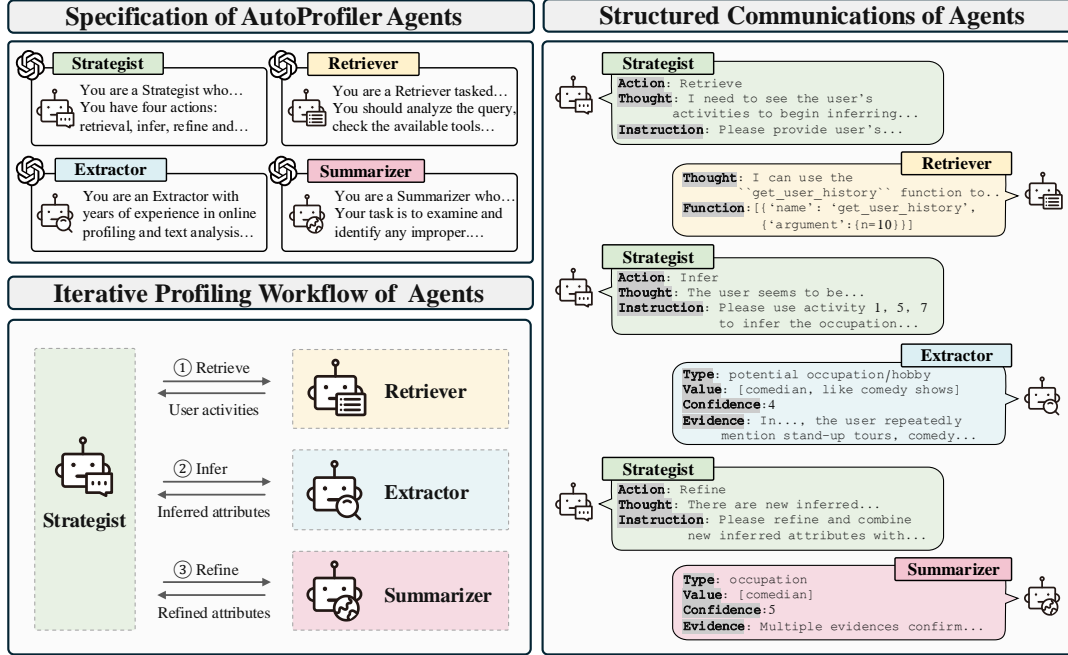


Figure 2: Illustration of the key profiling processes in AutoProfiler. Upper left: It employs four specialized agents to complete the task (full prompts are provided in Appendix D.1). Bottom left: Strategist coordinates other agents to sequentially retrieve, infer, and refine personal attributes. Right: Structured output of agents for efficient communication. Best viewed in color.

AutoProfiler consists of these four agents, each focusing on its own sub-tasks but collaborating systematically.

Specialization of Agents via Zero-shot Learning. To bypass the problem of providing suitable examples for agents, AutoProfiler employs zero-shot learning [29], enabling LLM agents to adapt without handcrafted demonstrations. Specifically, we provide detailed *descriptions* (e.g., defining what constitutes personal information) rather than specific demonstrations (e.g., showing how to infer a particular attribute from a given text). This approach enables agents to understand task objectives based on these descriptions, allowing them to interpret and perform tasks autonomously.

As illustrated in the upper left part of Figure 2, each agent is initialized with specialized instructions and tools tailored to its specific task (see Appendix D.1 for detailed prompts). Specifically, Strategist is instructed to plan the next steps based on the available user activities and inferred attributes; Retriever is guided in using API functions to collect user activities; Extractor is provided with criteria for identifying personal attributes; and Summarizer is assigned to verify and refine inferred attributes by checking for inconsistencies, ambiguities, inaccuracies, and duplicates.

Workflow across Agents. We design an *iterative* workflow that enables agents to profile users incrementally, processing one batch of activities per inference iteration. (The batch size is set to 10 activities to accommodate the context windows of different LLMs). The workflow is illustrated in the bottom left of Figure 2, with each action numbered for clarity. Each iteration begins with Strategist determining the next action. If no user activities have been collected, Strategist instructs Retriever to collect a batch of the user's new activities (process ① in Figure 2). Once these activities are

retrieved, Strategist evaluates whether they contain sufficient information for Extractor to analyze. If more information is needed, Strategist instructs Retriever to continue gathering additional activities. When enough information is available, Extractor proceeds to infer relevant personal information from the collected activities (process ② in Figure 2). The inferred information is then sent to Summarizer, who consolidates it with the existing profile and refines it by resolving inconsistencies, ambiguities, inaccuracies, and duplicates (process ③ in Figure 2). Finally, the refined profile is returned to Strategist, which initiates the next round of inference if necessary. This iterative process continues until Strategist issues a finish command, indicating that no further information can be extracted and no additional activities remain for analysis.

Communications between Agents. To facilitate effective communication among agents, we require agents to produce structured outputs (i.e., JSON format) rather than using natural language. This approach allows for better organization of information and more efficient exchanges. We establish a schema and format tailored to each agent's role, ensuring that the necessary outputs are clearly defined and consistent. As depicted on the right side of Figure 2, Strategist produces three key outputs: the action to take, the rationale for this action, and corresponding instructions. There are four possible actions: retrieve, infer, refine, and finish. The retrieve action directs Retriever to gather additional user activities from pseudonymous platforms. The infer action prompts Extractor to infer personal information based on the given context, while the refine action instructs Summarizer to re-examine inferred attributes to ensure accuracy and reliability. Notably, instead of directly sending messages to other agents, Strategist selects an action that defines the next step. This approach simplifies the workflow by focusing

Strategist on specific actions rather than requiring it to manage the entire network of agents and their roles. This makes it easier for LLMs to understand and execute each task without needing comprehensive awareness of the full agent environment.

Structured outputs are also implemented for Extractor and Summarizer. Each inferred piece of information is formatted in JSON with four attributes: type, value, confidence, and evidence. Confidence scores are required to range from 1 to 5, with higher scores indicating greater confidence in the inference. To account for potential uncertainty in inferences, we allow Extractor to suggest up to three possible values for each attribute. This structured approach enables Summarizer to efficiently validate the inferred information by assessing confidence levels and examining supporting evidence, thereby enhancing the reliability of the final profile. We also conduct experiments to demonstrate the effectiveness of using structured outputs for profiling, which are detailed in Appendix C.4.

Memory and Context management. Maintaining memory is essential for the profiling task, as Strategist needs to track task progression to determine appropriate next steps, and Summarizer relies on previously inferred attributes to resolve errors in newly inferred information. However, storing the entire history of agent interactions as memory is impractical due to the limited context window and the information overload challenges of LLMs [34].

To address this, we designed two memory mechanisms tailored to the agents' roles. For Extractor and Retriever, whose tasks are manageable with specific context provided by Strategist, we implement a *short-term* memory approach, retaining only the information necessary for a single inference loop. For Strategist and Summarizer, *long-term* memory is required to oversee and summarize the whole profiling process. Rather than storing all historical data, we use structured inferred attributes as a condensed representation of the history. This approach enables agents to retain critical information while minimizing context size and avoiding information overload. In our experiments, we find that this approach can efficiently handle a Reddit user with 1,530 comments without exceeding the context window limits of LLMs. We also evaluate the effectiveness of the proposed memory management in Appendix C.5.

Potential Enhancements of AutoProfiler. In our framework, Retriever is restricted to accessing only the user's activities. We recognize that equipping agents with additional tools and advanced LLM techniques could further improve the profiling effectiveness of AutoProfiler in practice. For example, Retriever could be enhanced with online search capabilities, such as Google Search, to update its knowledge and provide contextual information for Extractor's analysis. Another option is to allow Retriever to download all user activities for offline access and implement a retrieval-augmented generation (RAG) framework [33], which could provide supporting evidence for inference and help reduce LLM hallucinations [26]. Moreover, many online activities involve other modalities (e.g., photos), and state-of-the-art LLMs (e.g., GPT-4o) also support multi-modality. Combining textual data with other activities may yield a more comprehensive user profile. While these options offer promising directions for developing a more powerful automated profiling system, our findings in Section 5 indicate that AutoProfiler already performs impressively well and poses significant privacy risks. Therefore, we leave these enhancements for future work.

Discussion. We note that AutoProfiler intentionally does not incorporate a de-anonymization module for two reasons: (i) The goal of AutoProfiler is to infer detailed profiles, and the inferred attributes can be used to cause privacy breaches beyond de-anonymization (e.g., sensitive personal information leakage, discussed in Appendix E.1); and (ii) While it is technically feasible to integrate a de-anonymization module into AutoProfiler, doing so would raise significant ethical concerns, as we do not intend to cause large-scale real-world privacy breaches. Nevertheless, in Section 5.2 we present a case study demonstrating that the inferred attributes of AutoProfiler can indeed be used for de-anonymization.

5 Experiments

In this section, we present a series of comprehensive experiments to demonstrate the feasibility of automated profile inference as well as the effectiveness of AutoProfiler. We also provide a quantitative analysis of the inferred profiles and demonstrate how these profiles can violate privacy via de-anonymization. Specifically, we aim to answer the following sets of research questions:

- **RQ1:** How does AutoProfiler perform in automated profile inference on real-world pseudonymous platforms? What insights can be found from the inferred profiles, and what privacy risks do they pose to pseudonymous users?
- **RQ2:** How do the different components (*i.e.*, agents) and LLMs of AutoProfiler affect its performance? How about the efficiency and cost of AutoProfiler compared to that of human profilers?
- **RQ3:** How does AutoProfiler perform, compared with the state-of-the-art LLM-based method for sensitive attributes inference?

Evaluation Roadmap. One major challenge in evaluating profile inference is the lack of available datasets. To address this, we construct two real-world datasets (*i.e.*, Reddit and Twitter) and utilize one synthetic dataset to comprehensively evaluate AutoProfiler. (i) In Section 5.2, we use the Reddit dataset to demonstrate the feasibility of automated profile inference and the associated privacy risks. With real posts from active pseudonymous users, this dataset closely reflects the practical privacy risks of online pseudonymity. (ii) Section 5.3 presents results on the Twitter dataset, which evaluates the performance and efficiency of AutoProfiler, as well as the effectiveness of its components (*i.e.*, agents). Because the dataset comprises activities of verified public figures, the inferred information can be easily evaluated. (iii) In Section 5.4, we adopt the synthetic dataset [56], which includes human-curated ground truth, enabling a detailed comparison with the state-of-the-art approach. This dataset shows AutoProfiler's superiority over the baseline in terms of inference accuracy on PII attributes.

Each dataset contributes distinct insights to the evaluation process, together providing a thorough assessment of AutoProfiler's capabilities. The statistics for the used datasets are provided in Table 1, and the data construction and evaluation details are presented in the following subsections, with each dataset evaluated individually. The limitations of our evaluation are discussed in Section 7.

5.1 Experimental Setup

To evaluate AutoProfiler across different LLMs, we utilize the official LLM inference APIs provided by Alibaba, Anthropic, Google, and OpenAI. For Llama-3, we deploy the model locally on a server

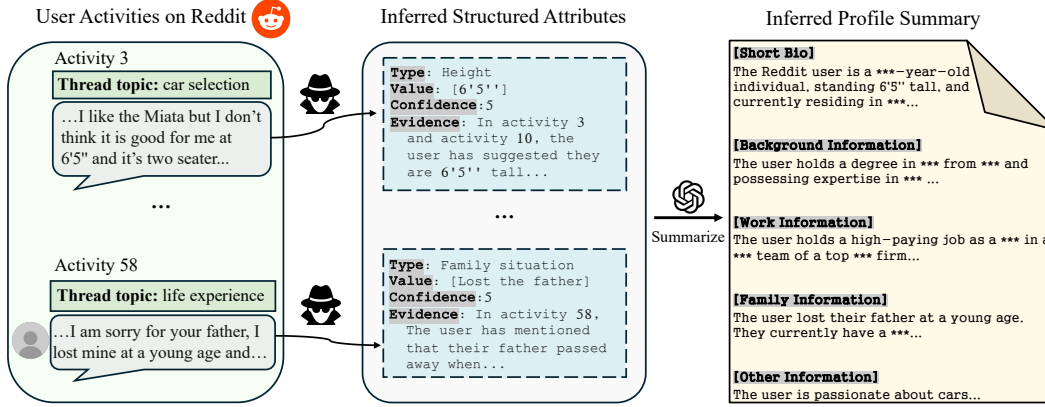


Figure 3: Demonstration of inferred attributes on Reddit dataset, the sensitive information is masked with “**”. AutoProfiler captures subtle clues (e.g., height) that users inadvertently reveal in seemingly insensitive contexts (e.g., car selection). We use GPT-4 to summarize these attributes for readability. Best viewed in color.**

Table 1: Summary of used datasets, including user numbers, activities per user, and words per activity.

Dataset	# Users	# Act. per User	# Words per Act.	Type
Reddit	250	857±230	42±40	Real-world
Twitter	100	126±118	27±14	Real-world
SynthPAI	300	26±20	19±15	Synthetic

Table 2: Inference accuracy (%) of AutoProfiler w.r.t generated confidence score. We randomly sampled 1,000 inferred attributes and manually evaluated their correctness.

Confidence of AutoProfiler	1	2	3	4	5
Inference accuracy	85%	88%	93%	100%	100%

equipped with four NVIDIA A100 GPUs. To enable online scraping functionalities, we provide Reddit and Twitter APIs for the Retriever agent. We employ AgentScope [18] framework to facilitate communication among multiple agents. Additional details regarding the LLMs and implementations can be found in Appendix A.

5.2 RQ1: Automated Profile Inference on Reddit

Roadmap. We evaluate AutoProfiler on Reddit, one of the largest pseudonymous online forums, to explore the effectiveness of AutoProfiler. We begin by outlining the data collection process and detailing the selection criteria for the user data. Then, we assess the reliability of the inferred attributes and provide illustrative examples. Next, we categorize these inferred attributes, conduct a quantitative analysis, and estimate the privacy risks associated with selected Reddit users based on these attributes. Finally, we present case studies demonstrating how inferred attributes could be leveraged to de-anonymize real Reddit users.

Dataset Construction. We use the following procedures to select users and their activities for constructing the dataset:

- (1) We follow [47] and select 438 popular subreddits where people are likely to discuss their personal matters.
- (2) For each subreddit, we extract the top 100 hot posts and record the users engaged in these threads to create a pool of candidates.
- (3) From this pool, we select the 250 most active users who participated across various subreddits, collecting their posts and comments from Jan 1, 2024, to May 31, 2024, ensuring that none of the used LLMs were trained on this data.

This process yielded a dataset of 250 Reddit users, with an average of 857 activities per user. Data collection was conducted through Reddit’s official API [27], which is publicly accessible and free to

use. It is worth noting that the users selected for our experiment do not represent general Reddit users: they are highly active users, therefore they are more vulnerable to profile inference. We include the complete subreddits used for data collection in Appendix B.

Hallucinations of Inferred Attributes. We use AutoProfiler with GPT-4 to infer personal attributes for selected Reddit users. To evaluate whether AutoProfiler produces inaccurate results due to LLM hallucinations [26], we randomly sampled 1,000 attributes and manually inspected their accuracy. Specifically, two of the paper’s authors independently performed the inspection, with inference accuracy determined by their full agreement. We then categorized accuracy based on the confidence scores generated by AutoProfiler, as shown in Table 2. The results indicate that the accuracy increases with higher confidence scores, with all attributes scoring above 3 aligning with human judgment. These findings suggest that AutoProfiler is reliable for inferring personal attributes, with confidence scores positively correlated with inference accuracy.

With the above observation, we *conservatively* filtered out attributes with confidence scores below 4 to obtain the final set of reliable inferred attributes for the following analysis. This yielded an average of 86 unique attributes per user, totaling 8,186 distinct attributes across the dataset.

Demonstrations of Inferred Attributes. We use GPT-4 to generate natural language summaries based on these inferred attributes to improve readability. Figure 3 illustrates the profile inference results for a Reddit user. In Activity 3, the user discusses car selection, specifically noting the limited space in a Miata and expressing concern about fitting comfortably—inadvertently hinting at his height. In another activity, the user expresses empathy for others by sharing the traumatic experience of losing the father at a young age. AutoProfiler captures these subtle implications and records them

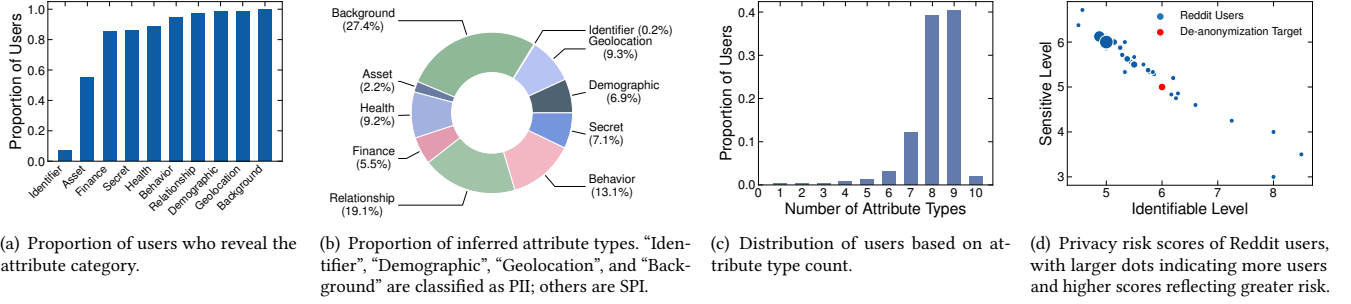


Figure 4: Analysis of the categorized attributes of Reddit users by category, count, and estimated privacy risks.

as inferred attributes about the user. These minor clues from online activities contribute to constructing a comprehensive user profile, capturing various facets of identity, including biological characteristics, education, employment, family background, and hobbies. We showcase more examples of inferred attributes in Appendix C.1.

Categorization of Inferred Attributes. We manually inspect the inferred attributes and classify them into two main categories for further analysis: Personally Identifiable Information (PII) and Sensitive Personal Information (SPI). (The rationale behind this categorization is provided in Appendix E.2.)

Personally Identifiable Information (PII). This category includes attributes that can be directly associated with an individual, encompassing the following four types:

- **Identifier:** Directly identifiable information, such as name, phone number, email address, *etc.*
 - **Demographic:** Attributes including age, gender, and ethnicity.
 - **Background:** Information about occupation, education, and achievements.
 - **Geographic:** Geo-information, such as birthplace and workplace.
- Sensitive Personal Information (SPI).* These attributes are sensitive but less likely to directly identify individuals. The pseudonymous nature of Reddit encourages users to share personal narratives, resulting in a considerable amount of SPI in posts and comments:
- **Health:** Physical or mental health conditions.
 - **Finance:** Financial records, wealth, and loan status.
 - **Relationship:** Information about family, marital status, and friends.
 - **Behavior:** Insights into routines, travel history, and commute, *etc.*
 - **Secrets:** Sensitive details, *e.g.*, sexual orientation and past traumas.
 - **Asset:** Real estate ownership, vehicle ownership, *etc.*

We used GPT-4 to classify all inferred attributes into ten predefined categories. Through manual inspection, we found that less than 1% of all inferred attributes were misclassified by GPT-4, which we then corrected to obtain the final categorized attributes (the prompts used and detailed inspection process are included in Appendix D.2). The results indicate that 43.8% of the inferred attributes are PII, while 56.2% fall under SPI.

Analysis of categorized Attributes. We present the statistics of categorized attributes in Figure 4. As shown in Figure 4(a) and Figure 4(b), fewer than 5% of pseudonymous users share identifiers online, and such information accounts for only 0.2% of all inferred attributes. This finding confirms that pseudonymous users generally avoid disclosing obvious personal information for their anonymity.

However, many pseudonymous users still reveal a significant amount of PII, such as their background and geographic locations, which could be used to infer a user’s identity. Additionally, we find

that Reddit users are willing to discuss sensitive topics, such as family matters and secrets. Although these conversations may not directly reveal an individual’s identity, they involve deeply personal content that can lead to unintended exposure.

What’s more concerning is that most users do not restrict themselves to discussing a single type of attribute on Reddit; instead, they participate across a variety of topics. While each piece of information may seem harmless on its own, the cumulative effect of discussing career, location, hobbies, relationships, and health across different threads builds a comprehensive profile about the user. Figure 4(c) illustrates this situation, showing that most users discuss at least 7 distinct types of personal attributes.

Privacy Risks Estimation. To quantitatively assess privacy risks for selected Reddit users, we assign different sensitivity and identifiability scores (ranging from 1 to 10) to each attribute type, indicating how easily it could identify the user and how sensitive it is to the individual (detailed in Appendix C.2). We then calculate each user’s average sensitivity and identifiability scores according to their inferred attributes, as shown in Figure 4(d).

The figure illustrates a clear inverse relationship between sensitivity and identifiability. Users who share highly sensitive information generally avoid revealing identifiable details, while those more open about their identities tend to share less sensitive content. Additionally, the distribution suggests that most Reddit users feel comfortable discussing sensitive topics while keeping their identities undisclosed.

Case study: De-anonymization on Real Reddit Users. We use de-anonymization to further illustrate the privacy risks posed by malicious profile inference. First, we follow [8] and define de-anonymization as follows:

Definition 1 ((n, k) -Deanonymization). Let D_u represent the set of inferred attributes of a user u , let $\mathcal{D}^{\text{aux}} = \{D_i^{\text{aux}} | i = [N]\}$ be an auxiliary dataset containing N records of real individuals. A matching function f takes D_u and \mathcal{D}^{aux} as inputs and outputs the number of matching attributes. The user u can be (n, k) -deanonymized w.r.t \mathcal{D}^{aux} if the following condition holds:

$$\sum_{D_i^{\text{aux}} \in \mathcal{D}^{\text{aux}}} \mathbb{1}[f(D_u, D_i^{\text{aux}}) \geq n] \leq k, \quad (1)$$

where $\mathbb{1}[\cdot]$ is the indicator function, which returns 1 if the condition $f(D_u, D_i^{\text{aux}}) \geq n$ is satisfied and 0 otherwise.

The above definition can be interpreted as the linkability risk of target user u w.r.t. auxiliary dataset \mathcal{D}^{aux} . The risk of de-anonymization is high when a user’s record can be linked to a

small number of individuals (*i.e.*, a small k) while sharing many overlapping attributes (*i.e.*, a large n) with records in \mathcal{D}^{aux} .

To demonstrate the feasibility of de-anonymization from inferred profiles, we leveraged LinkedIn as the auxiliary dataset, as it contains over 1 billion users with detailed, publicly available profiles. Specifically, we utilized LinkedIn’s People Search feature [9] to search for individuals by applying filters such as location, company, and education. We selected 10 users with the highest identifiable scores from the dataset, and by using search filters corresponding to the inferred attributes, we successfully narrowed the ambiguous list of potential matches to fewer than five LinkedIn profiles per target user (*i.e.*, can be at least (3, 5)-deanonymization).

We find one user is particularly vulnerable (indicated by the red dot in Figure 4(d)), achieving (4, 1)-deanonymization: by using four inferred attributes (*i.e.*, work location, occupation, educational background, and gender), only one LinkedIn profile matched the target user’s inferred attributes. Furthermore, we find that other inferred attributes not used for de-anonymization (*i.e.*, age and company) also matched the information in the LinkedIn profile, which strengthens the confidence of the de-anonymization result.

The above study demonstrates that by using a user’s online activities and public information like LinkedIn, an attacker could gain a comprehensive understanding of the user and even de-anonymize them with high confidence, posing a significant privacy risk to online pseudonymity. In Appendix C.6, we present additional de-anonymization results to extend the above case study and further examine how the volume of a user’s online activity impacts the feasibility of de-anonymization.

5.3 RQ2: Performance and Efficiency on Twitter

Roadmap. In this section, we evaluate AutoProfiler on Twitter, one of the world’s largest social media platforms. (We use “Twitter” instead of “X” in this paper for clarity and familiarity.) Unlike Reddit, many Twitter users use their real names for online interaction, presenting both opportunity and challenge to assess the performance of AutoProfiler. We begin by detailing the dataset preparation process and the text anonymization steps. Next, we outline the evaluation criteria, followed by a comprehensive assessment of AutoProfiler, including the impact of different LLM backbones and the contributions of each agent. Finally, we discuss the efficiency and cost of the proposed method.

Dataset Construction. Using Twitter’s official API [10], we selected the top 100 most-followed individual accounts, resulting in a dataset where all accounts are officially verified and the user’s name is known. To avoid data contamination in the LLMs, we only include their activity (*i.e.*, tweets) from Jan 1, 2024, to May 31, 2024. This dataset consists of 100 verified Twitter users, with an average of 126 tweets per user and 27 words per tweet. We provide the complete list of the selected Twitter users in Appendix B.

Tweet Anonymization. Tweets from verified Twitter users often contain specific information about themselves. For instance, singers may announce their live tours, and politicians may retweet news related to themselves. LLMs could potentially identify specific users by simply linking frequently mentioned names in tweets. To force LLMs to infer information from semantic cues rather than direct

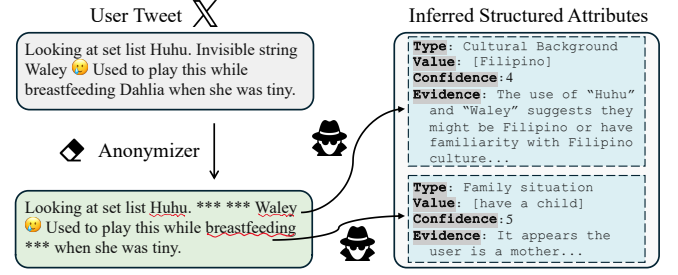


Figure 5: Inferring personal attributes from anonymized Tweets. “Invisible string” and “Dahlia” are masked as “****” as they refer to a song title and a person’s name. AutoProfiler still uncovers personal information like cultural background and family situation through subtle clues.

identifiers, we anonymize tweets by masking the identifiable entities. Specifically, we use state-of-the-art text anonymization tools from Azure [1] to detect and mask mentions of persons, locations, addresses, organizations, events, and numbers, replacing them with “****”, as shown in Figure 5. We use the anonymized Twitter dataset to assess the performance of AutoProfiler.

Evaluation Details. Directly evaluating the quality of inferred attributes is challenging due to the absence of ground truth. To address this, we design the following evaluation procedures for the Twitter dataset:

- (1) We first run AutoProfiler on the *anonymized* Twitter datasets to get the inferred attributes for each user.
- (2) Then we select attributes with high confidence scores (four or above) to construct a profile for each Twitter user and then instruct GPT-4 to use these inferred profiles to predict the user’s identity.
- (3) Finally, we measure the accuracy of these predictions as evaluation metrics, which is more stringent than just inferring attributes.

High-quality inferred attributes enable GPT-4 to more accurately identify users, resulting in higher identification accuracy. This approach is feasible because LLMs already possess prior knowledge of the well-known Twitter users used in our dataset.

Overall Performance. We evaluate AutoProfiler employing five state-of-the-art LLMs (*i.e.*, GPT-4, Claude-3, Gemini-1.5, Qwen-2, and Llama-3) as cores of agents. The results in Table 3 demonstrate that all LLMs achieve impressive identification accuracy, with at least 85% accuracy. GPT-4 performs the best, reaching an accuracy of 92% for well-known Twitter users. Notably, Llama-3, the locally deployed LLM model in our experiments, also performs well with an accuracy of 86%. This suggests that AutoProfiler could be used by an attacker to conduct large-scale automated profile inference at minimal cost and without centralized regulation (*e.g.*, advanced alignment strategies).

Ablation Study. To understand the impact of different agents on the final results, we conducted ablation studies to assess each agent’s contribution. Table 3 shows that incorporating additional agents beyond Extractor consistently enhances identification accuracy across all LLMs. Specifically, Strategist agent aids in identifying relevant types of personal information within tweets and helps design the inference strategy; Retriever effectively handles noisy or lengthy tweets; and Summarizer ensures the reliability of inferred attributes. By contrast, using a single LLM (*e.g.*, Extractor) yields

Table 3: Evaluation of AutoProfiler on the Twitter dataset. Identification accuracy is used to assess the performance and we use the GPT-4 to calculate the average token numbers and price per user. “✓” indicates the addition of a specific component.

Components of AutoProfiler				Performance (%)					Cost		
Extractor	Strategist	Retriever	Summarizer	GPT-4	Claude-3	Gemini-1.5	Qwen-2	Llama-3	# Input tokens	# Output tokens	Price (USD)
✓	✗	✗	✗	72±2	70±3	74±3	60±5	60±4	38,843	2,689	\$0.23
✓	✓	✗	✗	77±1	74±2	76±2	64±3	61±2	44,681	3,194	\$0.27
✓	✓	✓	✗	84±3	79±2	83±2	70±2	69±3	58,604	5,018	\$0.37
✓	✓	✗	✓	86±2	80±3	82±1	69±3	70±2	52,056	5,571	\$0.34
✓	✓	✓	✓	92±1	87±1	90±2	85±3	86±2	90,755	9,003	\$0.59

the lowest performance. Notably, weaker LLMs, such as Qwen-2 and Llama-3, derive greater benefit from these additional agents, showing substantial performance improvements when multiple agents are employed to complete the task.

Efficiency/Cost Evaluation. To evaluate the efficiency of AutoProfiler, we measure the input and output tokens required for LLMs throughout the entire inference process and estimate the cost based on GPT-4 pricing. As shown in Table 3, adding more agents to our framework increases communication costs (*i.e.*, input/output tokens) and overall expenses. Our experiments indicate that GPT-4 completes in roughly half a minute with OpenAI’s Batch service, whereas a human requires about an hour on average (including actions such as clicking, note-taking, and online information searches). Consequently, AutoProfiler achieves a 120× speed improvement and a 50× cost reduction. We note that this estimation is highly conservative; attackers could achieve better efficiency and cost-effectiveness by employing cheaper, faster, and more advanced models (*e.g.*, GPT-4o) or leveraging local LLMs (*e.g.*, Llama-3) for profiling. Detailed calculations can be found in Appendix C.3.

Discussion. We acknowledge that using verified Twitter users for evaluation is not without its limitations. However, to the best of our knowledge, there is currently no public dataset that provides grounded profiles for non-famous users. Therefore, we employ text anonymization and stricter evaluation criteria to analyze the efficiency and ablation of AutoProfiler. We also discuss the considerations behind the design of the evaluation in Appendix E.3.

5.4 RQ3: Comparison with the State-of-the-Art

Roadmap. To compare AutoProfiler with the state-of-the-art LLM-based inference method, we use SynthPAI [56], a recently proposed synthetic dataset containing human-labeled comments with predefined PII attributes. We begin by describing the dataset creation process and outlining the differences between AutoProfiler and the baseline. Next, we present a detailed performance comparison by examining prediction accuracy across different PII attributes. We then evaluate calibration accuracy, taking into account different levels of attribute hardness and certainty. Finally, we assess the robustness of AutoProfiler against incorrect information in activities.

Dataset Description. The SynthPAI dataset is constructed in three steps: (i) Creating diverse synthetic user profiles and initializing LLM agents with these profiles; (ii) Generating comments by enabling interactions between agents; (iii) Labeling the generated comments with predefined PII attributes, assisted by an LLM. The resulting dataset comprises over 7,800 comments from 300 synthetic users, totaling 700 ground-truth attributes. Each user is manually labeled with a subset of eight predefined PII attributes inferred

from their comments: Age, Sex, Education, Income Level, Relationship Status, Place of Birth, Location, and Occupation. Additionally, each attribute is annotated with *hardness* and *certainty* levels on a scale from 1 to 5, where higher values indicate greater difficulty and higher confidence, respectively. It is worth mentioning that SynthPAI intentionally removes explicit clues from comments, requiring LLMs to infer attributes based on subtle indicators such as dialect and cultural implications. However, as shown in Figure 3, in real-world scenarios, individuals often mention their traits more explicitly, resulting in more accurate and detailed profile inference.

Baseline and Evaluation Metric. Free text inference (FTI) [47] is the state-of-the-art LLM-based approach for PII inference, which is specifically tailored for use with the SynthPAI dataset. Specifically, it formulates the attribute inference task as a classification task: it directly feeds all user text to a single LLM instance and instructs it to select the most appropriate value for each PII attribute. This inference process is conducted in a single round, generating predictions for all attributes simultaneously. While AutoProfiler is designed as a general framework for automated profile inference, FTI offers an opportunity for performance comparisons in a less realistic setting. To ensure a fair comparison, we instruct AutoProfiler to predict the same designated attributes as FTI. If either method produces inferred values outside the predefined categories, we re-run both methods until the predictions align with these categories. We evaluate performance based on prediction accuracy.


Overall Performance. We used five LLMs as the backbone of AutoProfiler and FTI to compare the performance. As shown in Table 4, AutoProfiler outperforms FTI across all attributes and LLM backbones. We attribute this improvement to AutoProfiler’s iterative inference workflow: by analyzing and inferring smaller portions of text sequentially, LLMs can filter out noisy information and more effectively capture nuanced implications. In addition, AutoProfiler re-evaluates inferred attributes across different batches of user text, promoting consistency and enhancing the reliability of inferences. We note that the prediction accuracy for income is relatively low, even for AutoProfiler. We think this may be because the dataset deliberately omits any explicit income amounts on texts. Without these specific details, inferring income becomes inherently challenging and uncertain, even for human evaluators.

Calibration Accuracy. We demonstrate that the predictions of AutoProfiler align closely with human labels. To assess this alignment, we utilize calibration accuracy, which measures the prediction accuracy of AutoProfiler across different levels of inference hardness and certainty. We define calibration accuracy as:

Definition 2 (Calibration Accuracy). *Let $T = \{t_1, \dots, t_m\}$ be the set of attributes with ground truth values, let $\hat{T} = \{\hat{t}_1, \dots, \hat{t}_m\}$ be the set of inferred attributes predicted by AutoProfiler. A set $C_l \subseteq [m]$*

Table 4: PII prediction accuracy (%) on the SynthPAI dataset. AutoProfiler outperforms the baseline for all attributes and LLMs.

LLM	Method	AGE	EDU	INC	LOC	OCC	POB	REL	SEX
GPT-4	FTI	69.4	73.0	66.7	80.0	73.9	88.0	79.2	92.8
	AutoProfiler	80.6	81.0	75.6	88.1	95.4	92.0	89.6	93.7
Claude-3	FTI	47.2	69.0	64.5	71.2	75.4	78.0	86.5	91.9
	AutoProfiler	75.0	75.5	68.9	92.5	90.8	84.0	87.5	92.8
Gemini-1.5	FTI	66.7	53.5	51.1	66.3	65.7	84.0	78.1	76.6
	AutoProfiler	77.8	71.0	71.1	83.1	88.6	88.0	79.2	85.6
Qwen-2	FTI	50.0	59.0	40.0	76.9	71.1	80.0	72.9	88.3
	AutoProfiler	75.0	62.0	52.0	86.8	86.9	84.0	84.4	90.1
Llama-3	FTI	69.4	73.0	46.7	80.6	72.9	84.0	72.9	82.0
	AutoProfiler	75.0	75.0	50.0	81.3	85.5	92.0	77.1	84.7

**Figure 6: Calibration accuracy (Cal-Acc) of AutoProfiler on the SynthPAI dataset. Hardness and certainty scores are labeled by humans, and confidence is generated by AutoProfiler during prediction. Higher scores indicate greater difficulty, certainty, or confidence, respectively. The inferences made by humans and AutoProfiler are generally well-aligned.****Table 5: PII prediction accuracy (%) on original and noisy SynthPAI datasets. The performance remains stable despite some incorrect information in the activities.**

	AGE	EDU	INC	LOC	OCC	POB	REL	SEX
Original dataset	80.6	81.0	75.6	88.1	95.4	92.0	89.6	93.7
Noisy dataset	79.4	81.0	74.5	87.7	94.8	90.6	88.2	93.7

consists of attributes that belong to this specific annotated type l (e.g., hardness=3). The calibration accuracy for type l is defined as:

$$\text{Calibration Accuracy} = \frac{\sum_{j \in C_l} \mathbb{1}[t_j = \hat{t}_j]}{|C_l|} \quad (2)$$

where $\mathbb{1}[\cdot]$ is the indicator function.

For a well-aligned automatic profiling system, calibration accuracy should be higher for attributes that are easier for humans to infer and associated with greater certainty.

We present the calibration accuracy across each hardness and certainty level in Figure 6(a) and Figure 6(b), respectively. As shown, we observe a decrease in accuracy as hardness scores increase, indicating that models and human labelers generally agree on which examples are more challenging. A similar trend is seen with certainty levels: the prediction accuracy of AutoProfiler increases as human certainty scores rise.

We also examine the reliability of the confidence scores assigned by AutoProfiler. Specifically, AutoProfiler assigns a confidence score from 1 to 5 for each predicted attribute, with higher

scores indicating greater confidence in the prediction. Figure 6(c) shows that accuracy increases with confidence scores, suggesting that AutoProfiler’s confidence scores are a reliable indicator of prediction quality. Although the accuracy for confidence score 2 is 100%, this result may be due to the limited number of attributes assigned to this confidence score (10 out of 700).

Profiling with Noisy Activities. AutoProfiler introduces the Summarizer agent to address and refine incorrectly inferred attributes from noisy activities. We now evaluate its effectiveness through the following experiment. Specifically, for the synthetic dataset, we intentionally replaced 10% of a target user’s comments with randomly selected comments from other users, while keeping the ground truth (user profile) unchanged. We then used AutoProfiler to infer attributes from this noisy dataset. A performance comparison between the noisy and original datasets using GPT-4 is shown in Table 5. The results reveal only minimal degradation in performance, demonstrating that AutoProfiler effectively manages inconsistent or incorrect information in users’ activities.

6 Mitigation Strategies

In this section, we discuss potential mitigation strategies from the perspectives of users, pseudonymous online platforms, LLM providers, and privacy-enhancing technologies.

User-Side Mitigation. As the privacy threat discussed arises from user-generated activities, we advocate for increasing public awareness about the potential vulnerabilities of online pseudonymity. It is essential for individuals to understand these risks and exercise

Table 6: Comparing the identification accuracy (%) of AutoProfiler on raw and anonymized Twitter dataset.

	GPT-4	Claude-3	Gemini-1.5	Qwen-2	Llama-3
Raw Dataset	98%	93%	94%	92%	90%
Anonymized Dataset	92%	87%	90%	85%	86%

caution in online interactions. We also explore technical solutions to mitigate these threats, as outlined below.

A common approach to protecting sensitive information in text is to use text anonymizers [35]. For example, entity recognition tools [21] can be used to identify PII within the text, which can then be masked before training or publishing [35]. However, this approach is limited in preventing this threat for two main reasons:

- *Ineffectiveness of existing anonymizers.* We find that state-of-the-art text anonymizers are ineffective in preventing LLM-based profile inference. To illustrate this, we compare the raw Twitter dataset with its anonymized version, processed by Azure anonymizer [1], and evaluate identification accuracy across both versions. As shown in Table 6, while anonymization resulted in a slight decrease in accuracy, the overall accuracy remains significantly high. This is because AutoProfiler can infer personal information through contextual clues (see Figure 5)), whereas current anonymization tools focus on masking word-level sensitive information. This observation aligns with previous studies [47], which suggest that text anonymizers are insufficient for automated profile inference.

- *Infeasibility of anonymization for online activities.* Since AutoProfiler leverages publicly accessible user activities for profiling, text anonymization is often not an option in this context. Anonymizing users' posts may negatively impact user experience, restrict expressiveness, or even alter the original meaning.

Another way to address this threat is to develop detection tools that inform and alert users about their privacy leakage levels based on their online activities. Unfortunately, to the best of our knowledge, no such tool currently exists. We believe that creating such a tool could help individuals recognize the potential privacy risks of online pseudonymity and reduce personal information exposure. We leave this as a direction for future work.

Platform-Side Mitigation. We advocate two strategies for pseudonymous platforms to protect users' privacy against such threats. First, platforms could allow users to manage who can view their activities and adopt different pseudonyms to obscure their online personas. For instance, Reddit could offer users options to control post visibility or automatically hide older activities to enhance pseudonymity. Second, platforms may impose restrictions on API usage to prevent misuse. For example, setting limits on the number of retrievable activities would make it more difficult for attackers to build detailed profiles based on limited data.

LLM-Side Alignment. LLM alignment [4] is an active area of research on ensuring LLMs' outputs aligned with human values. However, we find that current LLMs are not effectively aligned against the privacy-invasive prompts used in AutoProfiler. Table 7 presents the average detection rate for unsafe prompts. Across all providers, we observe that most LLMs fail to identify malicious usage, with only a small percentage of requests flagged as unsafe by Google Gemini and Anthropic Claude. Additionally, even when

these prompts are detected as unsafe, users can still receive responses from the LLMs. We believe that more effective alignment methods are essential to help mitigate this privacy risk.

Privacy-enhancing Technologies. We find that existing privacy-enhancing technologies, such as k-anonymity [50] and differential privacy [15], are challenging to apply to the threat discussed in this paper. One reason is that the privacy risk stems from user-generated content, making it challenging to protect the sensitive information contained in the content before publishing. In addition, most existing privacy-enhancing methods require trade-offs that limit data utility, which is impractical for online communication, as it may impair user experience and lead to misunderstandings. We advocate for the privacy research community to develop new privacy-enhancing technologies to address this new threat.

7 Discussion & Ethical Considerations

Limitation of Evaluated Datasets. We note that the datasets used in our experiments have limitations for evaluation. For the Reddit dataset, we do not have the ground truth about the inferred attributes, thus it is challenging to directly assess the performance. The Twitter dataset, composed of Tweets from verified users, may not accurately reflect the online behavior of the general population. For the SynthPAI dataset, as previously discussed, there are significant distributional differences between synthetic data and real-world textual activities, and only ground truth for PII attributes is provided. While each dataset has its own limitations, the combined use of all three provides a comprehensive assessment of the performance of the proposed method. Nevertheless, there are no publicly available datasets tailored to evaluate this emerging threat.

Potential Use Cases of AutoProfiler. Although AutoProfiler is primarily proposed to explore the privacy risks of online pseudonymity, it can also be applied to other scenarios:

- *Privacy Risk Detection Tools.* AutoProfiler could be used as a privacy assessment tool to alert users to the privacy risks associated with their online activities. Additionally, it could serve as a defense against the threat identified in this paper by warning users of potential privacy risks before they share information online.

- *Criminal profiling.* Criminal profiling aims to identify the personality and behavioral characteristics of an offender, typically requiring the expertise of highly trained specialists [6]. We believe that AutoProfiler could be a valuable tool to support criminal profilers in efficiently capturing relevant traits of offenders, enhancing the effectiveness of criminal investigations.

Ethics Statement. This research was approved by the ethics committee of the authors' institution. All data used in this paper were obtained through official APIs and fully complied with the platform's regulations. No real human subjects were involved in our experiments. We disclosed our findings to major LLM providers, including Alibaba, Anthropic, Google, Meta, and OpenAI in December 2024. We have also notified Reddit and X about the potential de-anonymization risks for their users. We recognize that the results presented in this paper may raise concerns about privacy rights, especially given that current mitigation strategies are insufficient to fully address the threat. However, these actions were already possible before this research, and we believe that raising awareness is a crucial first step toward mitigating broader privacy risks.

Table 7: Percentage of unsafe requests detected by LLMs.

	GPT-4	Claude-3	Gemini-1.5	Qwen-2	Llama-3
Detection Ratio	0%	2.3%	8.4%	0%	0%

Datasets and Codes. Although AutoProfiler uses only publicly available data, we find that the inferred attributes of Reddit users are too sensitive to disclose. Additionally, Twitter’s API policy prohibits the republishing of tweets, even if they are publicly accessible. To support reproducibility, we will provide detailed descriptions of the data collection process to obtain our experimental datasets. We will also release the code² needed to reproduce results on the SynthPAI dataset, with modifications to prevent direct use for online profiling. We also note that the examples shown in Figure 2 and Figure 3 are synthetic to protect users’ privacy. These examples have been carefully crafted to ensure their core content closely reflects real samples without misleading readers.

8 Conclusion & Future Work

In this paper, we introduce a new privacy threat that LLMs pose to online pseudonymity called automated profile inference. Specifically, we design an LLM-based multi-agent framework called AutoProfiler, which can automatically scrape and extract sensitive personal attributes from publicly visible user activities on pseudonymous platforms. Experimental results on two real-world datasets and one synthetic dataset show that AutoProfiler is both effective and efficient and can be easily deployed on a web-scale. We reveal that the inferred attributes can lead to severe privacy breaches, such as de-anonymization. In addition, we find that AutoProfiler can capture subtle clues and implications from online activities and uncover deeply private information beyond PII, which may be exploited to pose even more severe privacy risks.

Moving forward, there are numerous rich areas for future work in privacy risks of online pseudonymity. In particular, the automated profile inference of posting photos and videos remains unexplored. Moreover, our work highlights challenges in mitigating this privacy threat and advocates for greater public awareness. We hope to tackle these and other challenges in future work.

References

- [1] Aahill. 2024. What is Azure AI Language? *Azure AI services* (2024). <https://learn.microsoft.com/en-us/azure/ai-services/language-service/overview>
- [2] Accountability Act. 1996. Health insurance portability and accountability act of 1996. *Public law* 104 (1996), 191.
- [3] Anthropic. 2024. Introducing the next generation of Claude. (2024). <https://www.anthropic.com/news/claude-3-family>
- [4] Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, Carol Chen, Catherine Olsson, Christopher Olah, Danny Hernandez, Dawn Drain, Deep Ganguli, Dustin Li, Eli Tran-Johnson, Ethan Perez, Jamie Kerr, Jared Mueller, Jeffrey Ladish, Joshua Landau, Kamal Ndousse, Kamile Lukosiute, Liane Lovitt, Michael Sellitto, Nelson Elhage, Nicholas Schiefer, Noemi Mercado, Nova Das-Sarma, Robert Lasenby, Robin Larson, Sam Ringer, Scott Johnston, Shauna Kravec, Sheer El Showk, Stanislav Fort, Tamera Lanham, Timothy Telleen-Lawton, Tom Conerly, Tom Henighan, Tristan Hume, Samuel R. Bowman, Zac Hatfield-Dodds, Ben Mann, Dario Amodei, Nicholas Joseph, Sam McCandlish, Tom Brown, and Jared Kaplan. 2022. Constitutional AI: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073* (2022).
- [5] Daniil A Boiko, Robert MacKnight, and Gabe Gomes. 2023. Emergent autonomous scientific research capabilities of large language models. *arXiv preprint arXiv:2304.05332* (2023).
- [6] Firearms Bureau of Alcohol, Tobacco and Explosives. [n. d.]. Criminal Profilers. [n. d.]. <https://www.atf.gov/careers/criminal-profilers>
- [7] Nicholas Carlini, Florian Tramèr, Eric Wallace, Matthew Jagielski, Ariel Herbert-Voss, Katherine Lee, Adam Roberts, Tom B. Brown, Dawn Song, Úlfar Erlingsson, Alina Oprea, and Colin Raffel. 2021. Extracting training data from large language models. In *30th USENIX Security Symposium*. 2633–2650.
- [8] Shuchi Chawla, Cynthia Dwork, Frank McSherry, Adam Smith, and Hoeteck Wee. 2005. Toward privacy in public databases. In *Theory of Cryptography: Second Theory of Cryptography Conference*. 363–385.
- [9] LinkedIn Corp. 2025. *LinkedIn search*. <https://www.linkedin.com/search/results/people>
- [10] X Corp. 2025. *X API*. <https://developer.x.com/en/docs/x-api>
- [11] Qingxiu Dong, Lei Li, Damai Dai, Ce Zheng, Zhiyong Wu, Baobao Chang, Xu Sun, Jingjing Xu, and Zhifang Su. 2022. A survey on in-context learning. *arXiv preprint arXiv:2301.00234* (2022).
- [12] David M Douglas. 2016. Doxing: a conceptual analysis. *Ethics and information technology* 18, 3 (2016), 199–210.
- [13] John E Douglas, Robert K Ressler, Ann W Burgess, and Carol R Hartman. 1986. Criminal profiling from crime scene analysis. *Behavioral Sciences & the Law* 4, 4 (1986), 401–421.
- [14] Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783* (2024).
- [15] Cynthia Dwork. 2006. Differential Privacy. In *Automata, Languages and Programming*. 1–12.
- [16] Dominique Estival, Tanja Gaustad, Son Bao Pham, Will Radford, and Ben Hutchinson. 2007. Author profiling for English emails. In *Proceedings of the 10th Conference of the Pacific Association for Computational Linguistics*, Vol. 263. 272.
- [17] Richard Fang, Rohan Bindu, Akul Gupta, Qiusi Zhan, and Daniel Kang. 2024. LLM agents can autonomously hack websites. *arXiv preprint arXiv:2402.06664* (2024).
- [18] Dawei Gao, Zitao Li, Weirui Kuang, Xuchen Pan, Daoyuan Chen, Zhijian Ma, Bingchen Qian, Liuyi Yao, Lin Zhu, Chen Cheng, Hongzhu Shi, Yaliang Li, Bolin Ding, and Jingren Zhou. 2024. AgentScope: A Flexible yet Robust Multi-Agent Platform. *arXiv preprint arXiv:2402.14034* (2024).
- [19] Philippe Golle. 2006. Revisiting the uniqueness of simple demographics in the US population. In *Proceedings of the 5th ACM Workshop on Privacy in Electronic Society*. 77–80.
- [20] Neil Zhenqiang Gong and Bin Liu. 2018. Attribute inference attacks in online social networks. *ACM Transactions on Privacy and Security* 21, 1 (2018), 1–30.
- [21] Kishaloy Halder, Alan Akbik, Josip Krapac, and Roland Vollgraf. 2020. Task Aware Representation of Sentences for Generic Text Classification. In *28th International Conference on Computational Linguistics*. 3202–3213.
- [22] Saul Hansell. 2006. AOL removes search data on vast group of web users. *New York Times* (2006). <https://www.nytimes.com/2006/08/08/business/media/08aol.html>
- [23] Jia He, Mukund Rungta, David Koleczek, Arshdeep Sekhon, Franklin X Wang, and Sadiq Hasan. 2024. Does Prompt Formatting Have Any Impact on LLM Performance? *arXiv preprint arXiv:2411.10541* (2024).
- [24] Jim Hollan and Scott Stornetta. 1992. Beyond being there. In *Proceedings of the SIGCHI conference on Human factors in computing systems*. 119–125.
- [25] Sirui Hong, Mingchen Zhuge, Jonathan Chen, Xiwu Zheng, Yuheng Cheng, Jinlin Wang, Ceyao Zhang, Zili Wang, Steven Ka Shing Yau, Zijuan Lin, Liyang Zhou, Chenyu Ran, Lingfeng Xiao, Chenglin Wu, and Jürgen Schmidhuber. 2024. MetaGPT: Meta Programming for A Multi-Agent Collaborative Framework. In *The Twelfth International Conference on Learning Representations*.
- [26] Lei Huang, Weijiang Yu, Weitao Ma, Weihong Zhong, Zhangyin Feng, Haotian Wang, Qianglong Chen, Weihua Peng, Xiaocheng Feng, Bing Qin, and Ting Liu. 2023. A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. *arXiv preprint arXiv:2311.05232* (2023).
- [27] Reddit Inc. 2025. *Reddit API overview*. <https://www.reddit.com/dev/api/>
- [28] Bargav Jayaraman and David Evans. 2022. Are attribute inference attacks just imputation? In *Proceedings of the 2022 ACM SIGSAC Conference on Computer and Communications Security*. 1569–1582.
- [29] Takeshi Kojima, Shixiang Shane Gu, Machel Reid, Yutaka Matsuo, and Yusuke Iwasawa. 2022. Large language models are zero-shot reasoners. *Advances in neural information processing systems* 35 (2022), 22199–22213.
- [30] Michal Kosinski, David Stillwell, and Thore Graepel. 2013. Private traits and attributes are predictable from digital records of human behavior. *Proceedings of the national academy of sciences* 110, 15 (2013), 5802–5805.
- [31] Noam Lapidot-Lefler and Azy Barak. 2012. Effects of anonymity, invisibility, and lack of eye-contact on toxic online disinhibition. *Computers in human behavior* 28, 2 (2012), 434–443.
- [32] Marcus Law. 2023. Scam email cyber attacks increase after rise of ChatGPT. *Technology* (2023). <https://technologymagazine.com/articles/scam-email-cyber-attacks-increase-after-rise-of-chatgpt>
- [33] Patrick S. H. Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimír Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Advances in Neural Information Processing*

²<https://github.com/zealscott/AutoProfiler>

- Systems* 33 (2020), 9459–9474.
- [34] Nelson F Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. 2024. Lost in the Middle: How Language Models Use Long Contexts. *Transactions of the Association for Computational Linguistics* 11 (2024), 157–173.
 - [35] Nils Lukas, Ahmed Salem, Robert Sim, Shruti Tople, Lukas Wutschitz, and Santiago Zanella-Béguelin. 2023. Analyzing leakage of personally identifiable information in language models. In *IEEE Symposium on Security and Privacy*. 346–363.
 - [36] Shagufta Mehnaz, Sayanton V. Dibbo, Ehsanul Kabir, Ninghui Li, and Elisa Bertino. 2022. Are Your Sensitive Attributes Private? Novel Model Inversion Attribute Inference Attacks on Classification Models. In *31st USENIX Security Symposium*. 4579–4596.
 - [37] Niloofar Mireshghallah, Hyunwoo Kim, Xuhui Zhou, Yulia Tsvetkov, Maarten Sap, Reza Shokri, and Yejin Choi. 2024. Can LLMs Keep a Secret? Testing Privacy Implications of Language Models via Contextual Integrity Theory. In *The Twelfth International Conference on Learning Representations*.
 - [38] Arvind Narayanan and Vitaly Shmatikov. 2008. Robust de-anonymization of large sparse datasets. In *29th IEEE Symposium on Security and Privacy*. 111–125.
 - [39] State of California Legislature. 2018. California Consumer Privacy Act of 2018. *Public law* (2018).
 - [40] OpenAI. 2023. GPT-4 technical report. *arXiv preprint arXiv:2303.08774* (2023).
 - [41] Joon Sung Park, Joseph O'Brien, Carrie Jun Cai, Meredith Ringel Morris, Percy Liang, and Michael S Bernstein. 2023. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th annual acm symposium on user interface software and technology*. 1–22.
 - [42] Sundar Pichai and Demis Hassabis. 2024. Our next-generation model: Gemini 1.5. (2024). <https://blog.google/technology/ai/google-gemini-next-generation-model-february-2024/>
 - [43] Francisco Rangel, Paolo Rosso, Moshe Koppel, Efstathios Stamatatos, and Giacomo Inches. 2013. Overview of the author profiling task at PAN 2013. In *CLEF conference on multilingual and multimodal information access evaluation*. 352–365.
 - [44] Protection Regulation. 2016. Regulation (EU) 2016/679 of the European Parliament and of the Council. *Regulation (EU) 679* (2016).
 - [45] Thais Sardá, Simone Natale, Nikos Sotirakopoulos, and Mark Monaghan. 2019. Understanding online anonymity. *Media, Culture & Society* 41, 4 (2019), 557–564.
 - [46] Jasveer Singh. 2023. Pseudonymous platforms: The future of secure online communication. *The times of India* (2023). <https://timesofindia.indiatimes.com/blogs/voices/pseudonymous-platforms-the-future-of-secure-online-communication/>
 - [47] Robin Staab, Mark Vero, Mislav Balunović, and Martin Vechev. 2024. Beyond Memorization: Violating Privacy via Inference with Large Language Models. In *The Twelfth International Conference on Learning Representations*.
 - [48] Peter Steiner. 1979. On the Internet, nobody knows you're a dog. *The New Yorker* (1979).
 - [49] Latanya Sweeney. 1997. Weaving technology and policy together to maintain confidentiality. *The Journal of Law, Medicine & Ethics* 25, 2-3 (1997), 98–110.
 - [50] Latanya Sweeney. 2002. k-anonymity: A model for protecting privacy. *International journal of uncertainty, fuzziness and knowledge-based systems* 10, 05 (2002), 557–570.
 - [51] Saad Ullah, Mingji Han, Saurabh Pujar, Hammond Pearce, Ayse Coskun, and Gianluca Stringhini. 2024. LLMs Cannot Reliably Identify and Reason About Security Vulnerabilities (Yet?): A Comprehensive Evaluation, Framework, and Benchmarks. In *2024 IEEE Symposium on Security and Privacy*. 199–199.
 - [52] Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, Wayne Xin Zhao, Zhewei Wei, and Jirong Wen. 2024. A survey on large language model based autonomous agents. *Frontiers of Computer Science* 18, 6 (2024), 186345.
 - [53] Jason Wei, Yi Tay, Rishi Bommasani, Colin Raffel, Barret Zoph, Sebastian Borgeaud, Dani Yogatama, Maarten Bosma, Denny Zhou, Donald Metzler, Ed H. Chi, Tatsunori Hashimoto, Oriol Vinyals, Percy Liang, Jeff Dean, and William Fedus. 2022. Emergent abilities of large language models. *arXiv preprint arXiv:2206.07682* (2022).
 - [54] An Yang, Baosong Yang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Zhou, Chengpeng Li, Chengyuan Li, Dayiheng Liu, Fei Huang, Guanting Dong, Haoran Wei, Huan Lin, Jialong Tang, Jialin Wang, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin Ma, Jin Xu, Jingren Zhou, Jinze Bai, Jinzheng He, Junyang Lin, Kai Dang, Keming Lu, Keqin Chen, Kexin Yang, Mei Li, Mingfeng Xue, Na Ni, Pei Zhang, Peng Wang, Ru Peng, Rui Men, Ruize Gao, Runji Lin, Shijie Wang, Shuai Bai, Sinan Tan, Tianhang Zhu, Tianhao Li, Tianyu Liu, Wenbin Ge, Xiaodong Deng, Xiaohuan Zhou, Xingzhang Ren, Xinyu Zhang, Xipin Wei, Xuancheng Ren, Yang Fan, Yang Yao, Yichang Zhang, Yu Wan, Yunfei Chu, Yaqiong Liu, Zeyu Cui, Zhenru Zhang, and Zhihao Fan. 2024. Qwen2 Technical Report. *arXiv preprint arXiv:2407.10671* (2024).
 - [55] John Yang, Carlos E. Jimenez, Alexander Wettig, Kilian Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. 2024. SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering. *arXiv:2405.15793*
 - [56] Hanna Yukhymenko, Robin Staab, Mark Vero, and Martin T. Vechev. 2024. A Synthetic Dataset for Personal Attribute Inference. *Advances in Neural Information Processing Systems* 37 (2024).
 - [57] Elena Zheleva and Lise Getoor. 2009. To join or not to join: the illusion of privacy in social networks with mixed public and private user profiles. In *Proceedings of the 18th international conference on World wide web*. 531–540.

[illegible]

Figure 7: Complete SubReddit lists used for constructing the Reddit dataset.

We implement AutoProfiler in Python and use AgentScope [18] framework to facilitate communications among multiple agents. Specifically, we use the following LLMs for our experiments:

- We use the default hyperparameters of these models for inference. All models were provided the same prompts, provided in Appendix D.1. Additionally, we number each user activity chronologically to help agents retrieve and identify related activities for profile inference.

SubReddits Collections. We present a complete list of subreddits used to select the target Reddit users in Reddit dataset, as shown in Figure 7. The dataset is collected on June 28, 2024.

³<https://notcommon.com/most-followed/twitter>

Based on the ten types of inferred attributes defined in this paper, we assigned each a score from 1 to 10, reflecting both sensitivity and identifiability. This assignment follows a two-step process: First, each author independently assigns a unique score (ranging from 1 to 10) to each category. Then, the authors collaboratively review and refine these scores to reach a consensus. The scores are presented in Figure 10. We recommend that future work develop a

⁴<https://developer.x.com/en/docs/x-api>

⁵<https://www.tweepy.org/>

Sensitivity score	Identifiability score
Secrets: 10	Identifier: 10
Health: 9	Geographic: 9
Relationship: 8	Background: 8
Finance: 7	Demographic: 7
Behavior: 6	Asset: 6
Asset: 5	Relationship: 5
Background: 4	Behavior: 4
Geographic: 3	Finance: 3
Demographic: 2	Health: 2
Identifier: 1	Secrets: 1

Figure 10: Assigned sensitivity and identifiability scores for each inferred attribute type.

Table 8: PII prediction accuracy (%) on the SynthPAI dataset when using structured (JSON) or plain-text for agent communication.

	AGE	EDU	INC	LOC	OCC	POB	REL	SEX
Plain text	72.6	75.4	70.1	81.3	84.5	80.6	80.5	90.2
JSON	80.6	81.0	75.6	88.1	95.4	92.0	89.6	93.7

more systematic approach to assess the risks associated with these attributes.

C.3 Achievable Speedup

We present our calculations for the reported time (120×) and cost (50×) speedups achieved in profiling Twitter users. These values represent a comparison between a single human manually profiling a Twitter user and a single individual running our automated script. To protect user privacy, we did not use crowdsourcing for human labeling estimates; instead, the labeling was performed solely by the authors of this paper.

Our findings indicate that GPT-4, when run on OpenAI’s Batch service⁶, requires approximately 30 seconds to profile a user, whereas a human labeler requires around an hour, which includes actions such as clicking, note-taking, and online information searches. This results in a 120× time speedup with GPT-4. For the cost analysis, we assumed a standard rate of 30 USD per hour for human labeling, while the average cost for using GPT-4 is approximately 0.59 per user, based on OpenAI’s pricing⁷. This yields a cost reduction of around 50× when using GPT-4 for profiling.

It is also worth noting that the inference speed of LLMs is improving rapidly. Newer models, such as GPT-4o, offer faster inference speeds and lower costs, which may further increase the speed and cost advantages over human labeling.

C.4 Impact of Structured Communications between Agents

Recent research [23] has shown that formatting prompts in JSON leads to better and more stable performance compared to using plain text. Therefore, we conducted additional experiments to evaluate the effectiveness of using structured (JSON) outputs for agent communication in AutoProfiler. Specifically, we tested and compared

⁶<https://platform.openai.com/docs/guides/batch>

⁷<https://openai.com/api/pricing/>

Table 9: PII prediction accuracy (%) on the SynthPAI dataset with and without the memory management proposed in Section 4.2.

	AGE	EDU	INC	LOC	OCC	POB	REL	SEX
w/o memory	60.5	67.4	62.9	70.2	66.8	74.5	65.2	84.3
w/ memory	80.6	81.0	75.6	88.1	95.4	92.0	89.6	93.7

Table 10: De-anonymization results on the Reddit dataset. We use the size of the anonymity set as the metric. A smaller anonymity set size indicates greater vulnerability to de-anonymization.

# Comments per user	# < 300	# 300-400	# 400-500	# 500-600	# 600-700	# 700-800	# > 800
# attributes per user	28	44	48	51	60	85	99
Anonymity set size	88	73	61	42	32	21	10

Table 11: De-anonymization results on the synthetic dataset (i.e., SynthPAI). We use the top-1 and top-2 re-identification accuracy as the evaluation metric.

# Comments per user	# < 15	# 15-20	# 20-25	# > 25
Top-1 accuracy	0.66	0.69	0.73	0.88
Top-2 accuracy	0.83	0.92	0.95	0.94

the performance when Extractor and Summarizer communicated via free-text messages instead of JSON outputs in AutoProfiler, using GPT-4 on the synthetic dataset (i.e., SynthPAI). The results in Table 8 demonstrate that using JSON significantly improves the performance of profiling tasks across various attributes.

C.5 Impact of Memory Management in AutoProfiler

We now explore the impact of the proposed memory management for agents. Specifically, we remove the memory management and instead use the context windows of LLM agents to store all communication histories (clearing the history when the context window is exceeded), and compare their performance on the synthetic dataset (i.e., SynthPAI). The results in Figure 9 show a significant degradation in performance without memory management. This notable drop in performance highlights the importance of memory management for effective profiling in AutoProfiler.

C.6 Additional De-anonymization Results

Here we present additional experiments to further explore the feasibility of de-anonymization using inferred profiles. Specifically, we conducted the following two experiments:

- **De-anonymization on Reddit dataset.** We extended our case study in Section 5.2 to assess the de-anonymization risks associated with inferred attributes on the Reddit dataset. To do this, we divided the dataset into seven subgroups based on the number of comments per user, ranging from fewer than 300 to more

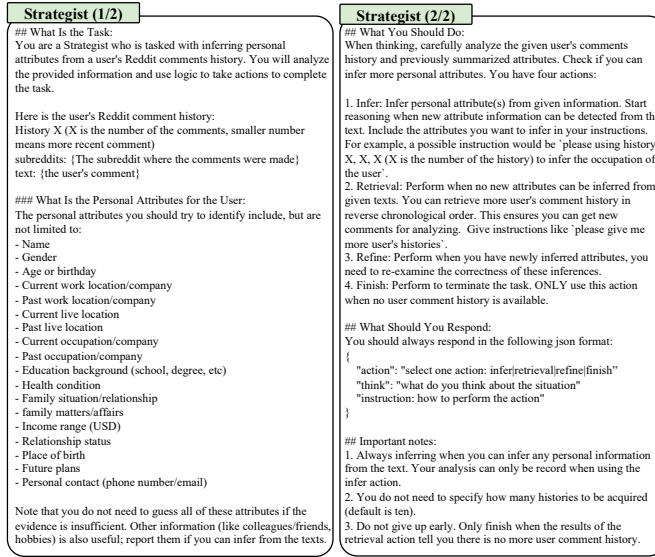


Figure 11: Complete prompts used for the Strategist agent.

than 800. For each subgroup, we randomly selected five users and manually searched for potential matches on LinkedIn using the same approach described in Section 5.2. We evaluated the de-anonymization by measuring the size of the anonymity set (i.e., the number of remaining potential matches on LinkedIn).

- **De-anonymization on synthetic dataset.** For the synthetic dataset (i.e., SynthPAI [56]), we used the ground-truth identity information to measure de-anonymization accuracy. Specifically, the synthetic dataset is generated by simulating conversations between LLM agents with predefined profiles (including eight PIIIs), and we can directly measure the de-anonymization rate based on the inferred attributes. We treated all users with profiles as anonymized candidates and, for each target user with inferred attributes, ranked all candidates using Hamming distance (candidates with identical profiles are ranked together). The top-k de-anonymization rate (accuracy) was then calculated as the metric.

Table 10 and Table 11 show the performance of the two experiments with respect to the number of comments posted by the user. These results share a similar trend: as AutoProfiler gains access to more of the user's activities, it is able to infer more personal attributes, thereby increasing the de-anonymization risk.

D Complete Prompts

D.1 Prompts for AutoProfiler Agents

Figure 11, 12, 13, and 14 present the complete set of prompts used for Strategist, Retriever, Extractor and Summarizer, respectively. These prompts define the roles of agents in conducting automated profiling tasks. Specifically, each prompt outlines the agents' responsibilities and requires them to produce structured outputs, facilitating effective communication between agents.

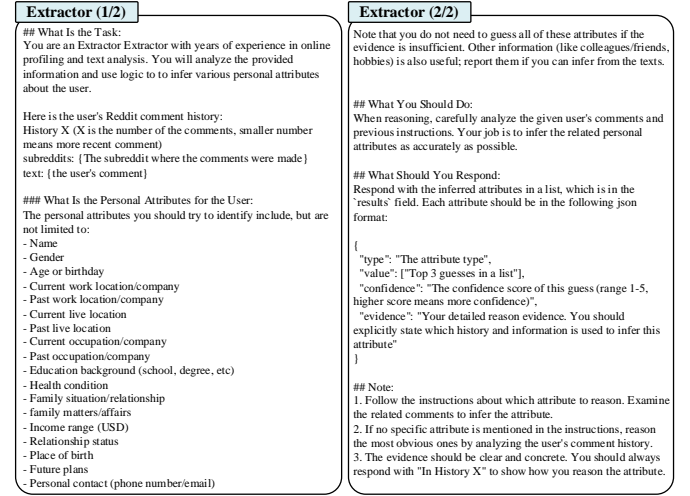


Figure 12: Complete prompts used for the Extractor agent.




Figure 13: Complete prompts used for the Retriever agent. The tool instructions are auto-generated by AgentScope [18] and are omitted here.

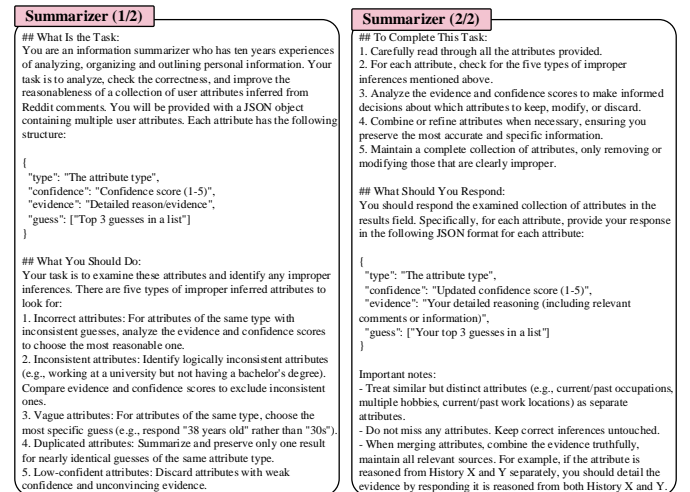
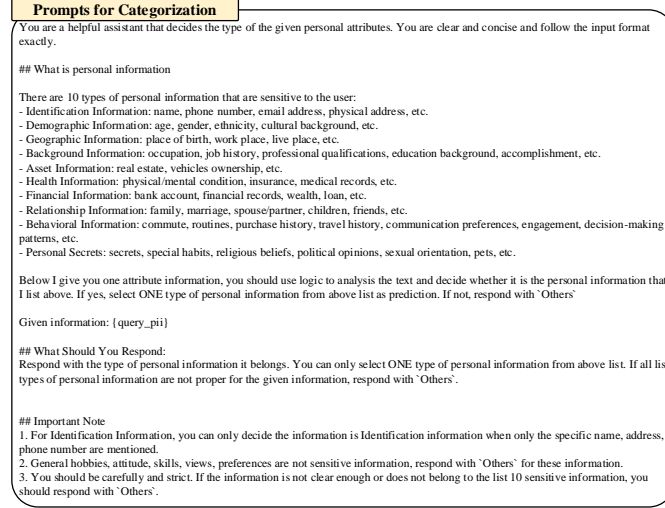


Figure 14: Complete prompts used for the Summarizer agent.

Table 12: Classification error of GPT-4, determined through manual inspection of all inferred attributes in the Reddit dataset.

	Identifier	Demographic	Background	Geographic	Health	Finance	Relationship	Behavior	Secrets	Asset
Error Rate	0%	0.01%	0.02%	0%	0.005%	0.11%	0.13%	0.18%	0.06%	0.003%

**Figure 15: Complete prompts used for categorizing inferred attributes.**

D.2 Prompts for Categorizing Inferred Attributes

Figure 15 presents the complete prompts used to classify all inferred attributes into ten predefined categories. We (two of the authors) manually inspected the classification outputs of all attributes, and the error rates for each category are shown in Table 12. As indicated, GPT-4 performs well in categorizing “Identifier” and “Geographic” information but exhibits slightly higher error rates in sensitive categories, such as “Finance” and “Behavior”.

E Additional Discussion

E.1 Privacy Risks of Sensitive Personal Information

Malicious attackers can exploit Sensitive Personal Information (SPI) inferred from AutoProfiler to carry out targeted attacks without needing to know the user’s identity. To illustrate the feasibility of this, we present two examples:

- **Health-Based Phishing.** If an attacker discovers a pseudonymous user has a specific health condition (e.g., diabetes), they could send tailored phishing messages impersonating a health-care provider. The victim, recognizing the relevance to their condition, is more likely to engage, inadvertently sharing credentials or payment details.
- **Relationship-based Impersonation.** Knowledge of a user’s secrets (e.g., undisclosed sexual orientation) or close relationships (e.g., a child’s nickname) allows attackers to craft believable

scams. For example, posing as a trusted friend via pseudonymous channels, the attacker could extort money or extract further sensitive details, leveraging emotional triggers without ever knowing the victim’s identity.

We refer to [12] for a detailed analysis of the consequences of exposing sensitive personal information.

E.2 Categorization of Inferred Attributes

Personally identifiable information (PII) can be a *direct* identifier when leakage of that data alone is sufficient to re-identify an individual, or a *quasi-identifier* when only an aggregation of many quasi-identifiers can reliably re-identify an individual [44, 47]. Information such as occupation and education is widely recognized as quasi-identifiers by existing research [35, 47], as they can contribute to re-identifying individuals when aggregated with other attributes. For example, as demonstrated in 5.2, backgrounds like occupation and education, when combined with other attributes (e.g., location), can significantly increase the likelihood of identifying a user when the auxiliary dataset is a professional platform like LinkedIn. We include both direct identifiers and quasi-identifiers in our PII categorization.

In our experiments, any inferred personal information that goes beyond PII is treated as sensitive personal information (SPI) and is manually divided into six categories. We acknowledge that this classification is by no means complete or perfect, and a systematic analysis of the SPI categorization is beyond the scope of this paper; we encourage future research to explore this further.

E.3 Evaluation for the Twitter Dataset

In our experiments, given the inferred attributes of the verified Twitter user, we use the identification accuracy of GPT-4 as the evaluation metric. Since LLMs already have some information about these verified users, a more straightforward approach would be to directly evaluate the correctness of the inferred information by prompting the LLMs to retrieve the ground truth. However, we do not choose this approach for two main reasons: (i) The online activities we used are more up-to-date than the LLMs’ training data, meaning that the inferred information and the LLM’s memorized information may not be coherent. For example, a CEO may have recently switched to another company, but the LLM may not be aware of this update, leading to errors in evaluation. (ii) It is well known that LLMs suffer from hallucination [26], meaning that the ground truth derived by prompting the LLMs may contain false information, thereby diminishing the reliability of the evaluation. Given the above considerations, we opt for a more rigorous evaluation approach by directly asking the LLMs to guess the identity of the user.